



ORBYS360

AI Knowledge Hub for GCCs

RESEARCH REPORT

THE GCC RECKONING

How AI Is Rewriting the Economics of Global Capability
Centers

Why half the GCC workforce will be redefined.

Orbys Research • 2026

SPONSORED BY



About Orbys360

The AI Knowledge Hub for Global Capability Centers.

Orbys360 is an AI-first knowledge platform built for the leaders running, modernising, and scaling Global Capability Centers. We work from a simple premise: GCCs are no longer support centers—they are becoming the AI backbone of the enterprises that own them. The window for that transition is short, the playbooks are scattered, and the decisions in front of GCC leadership are too consequential to be made on instinct.

We exist to close the gap between AI ambition and AI readiness. The Orbys360 platform brings together the resources GCC leaders need to move faster and with less risk:

- **Agent Marketplace** — a curated catalogue of production-ready AI agents organised by function and industry, with analyst profiles of the providers behind them.
- **Playbooks, Frameworks & Operating Models** — proven transformation methodologies, role taxonomies, and blueprints for AI-enabled operations.
- **AI Roadmaps & Implementation Templates** — strategic maturity pathways and project-ready resources from setup through scale.
- **AI Jobs for GCCs** — a curated board of AI-era GCC roles (agent operators, model validators, governance leads, prompt and integration specialists) paired with reskilling pathways that help leaders move existing talent into the work AI cannot do alone.
- **Research & Benchmarks** — data-driven comparative analysis, including the GCC workforce intelligence underpinning this report.
- **Expert Advisory & AI Pulse** — guidance from AI strategists, plus a curated daily briefing on the developments that change the calculus for GCC leadership.

If *The GCC Reckoning* reflects the shape of the problem, the Orbys360 platform is built to help leaders meet it—responsibly, and at scale.

Contents

About Orbys360	2
1 Executive Summary	7
2 The GCC Model: A Twenty-Year Retrospective	8
3 The Numbers: Aggregate Displacement Data	9
4 The Displacement Already Underway	9
4.1 The Corporate Record	10
4.2 The Industry Pattern	11
5 The Real Economics: Cost Arbitrage Made Visceral	12
5.1 The Math That Changes Boardrooms	12
5.2 The Team-Level View	13
6 Functional Threat Matrix	14
6.1 BPM — Business Process Management	15
6.2 IT Activities	17
6.3 ER&D — Engineering Research & Development	18
7 Role-Level Displacement Analysis	20
8 The Technology Driving Displacement	21
9 Sector Analysis: Financial Services GCCs	22
10 Sector Analysis: Tech GCCs	23
11 Sector Analysis: Healthcare and Life Sciences GCCs	24
12 Sector Analysis: Manufacturing GCCs	25
13 Sector Analysis: Retail GCCs	26
14 Sector Analysis: Energy GCCs	27
15 The Timeline: From Language Models to Agentic Systems	29
16 The Acceleration Curve: AI Capability vs. GCC Hiring Cycle	30
16.1 The Timing Mismatch	30
16.2 The Hiring Treadmill	31

17 The Structural Paradox	31
18 Enterprise HQ Sentiment: The Bloat Problem	33
18.1 The Three Patterns of GCC Bloat	33
19 The Seat Warmer Economy: How GCCs Built Their Own Structural Vulnerability	34
19.1 Defining the Seat Warmer: Role Archetype, Not Performance Verdict	34
19.2 Structural Origins: Five Mechanisms That Built the Archetype at Scale	36
19.3 The Hidden Cost Structure of Seat Warmer Economies	38
19.4 Structural Alignment with the AI Reckoning	40
19.5 The Path Out: From Seat Warmer Economy to Outcome Organisation	41
20 The Double Squeeze: AI From Below, Headquarters From Above	43
21 Geographic Impact: India	44
21.0.1 Function Composition and AI Risk	44
21.0.2 GCC Employment by City	44
21.0.3 Key Risk Factors	45
22 Geographic Impact: Eastern Europe	46
22.1 Poland	46
22.1.1 Function Composition and AI Risk	46
22.1.2 GCC Employment by City	46
22.1.3 Key Risk Factors	46
22.2 Romania	47
22.2.1 Function Composition and AI Risk	47
22.2.2 GCC Employment by City	47
22.2.3 Key Risk Factors	47
22.3 Czech Republic and Hungary	48
22.4 Bulgaria	48
22.5 Regional Pattern	48
23 Geographic Impact: South East Asia	49
23.1 The Philippines	49
23.1.1 Sector Composition and AI Risk	49
23.1.2 GCC Employment by Metro Area	49
23.1.3 Key Risk Factors	49
23.2 Vietnam	50
23.3 Indonesia	50
23.4 Malaysia	51
23.4.1 Function Composition and AI Risk	51
23.4.2 GCC Employment by Location	51
23.4.3 Key Risk Factors	51

24 Geographic Impact: Latin America	52
24.1 Mexico	52
24.2 Brazil	53
24.3 Colombia and Costa Rica	53
25 Geographic Impact: Middle East and Africa	53
25.1 Egypt	53
25.1.1 Function Composition and AI Risk	54
25.1.2 GCC Employment by Location	54
25.1.3 Key Risk Factors	54
25.2 Morocco	55
25.3 Rwanda and East Africa	55
26 The New Geographic Competitor: AI-Augmented Labour Markets	55
26.1 The Double Arbitrage Threat	56
27 The Labour Economics: Wage Compression, Hiring Freezes, and Shadow Attrition	57
28 The Human Cost	58
29 What Survives: The 50% and the Roles AI Cannot Replace	59
30 The New GCC Value Proposition: From Cost Arbitrage to Intelligence Hub	59
31 The Service Provider in the GCC Reckoning	61
31.1 The Service Provider's Own Reckoning	61
31.1.1 Service Provider AI Disruption Matrix	62
31.1.2 The Three Responses—and Their Outcomes	62
31.2 Hybridisation: From Competitor to Architecture Partner	64
31.2.1 Model One: Managed Capacity Transfer	64
31.2.2 Model Two: AI-Augmented Hub Co-Investment	64
31.2.3 Model Three: The Innovation Spine	65
31.3 Value Addition Through Managed Models	66
31.3.1 Seven Managed Model Plays	67
31.3.2 The Revenue Model Shift	67
31.3.3 The Financial Case	68
31.4 What GCC Clients Are Now Buying	69
31.5 Orbys360 Recommendation for Service Providers	69
32 Financial Modelling: The Cost of Inaction Versus the Cost of Transition	70
32.1 The Cost of Unmanaged Displacement	70
32.2 The Cost of Managed Transition	71
32.3 The Cost of Delay	71

33 Regulatory and Policy Environment	72
34 Strategic Recommendations for GCC Leadership	73
34.1 The Leadership Competency Gap	74
35 Strategic Recommendations for Enterprise Headquarters	74
36 Strategic Recommendations for National Governments	75
37 About This Research	76
38 Research Methodology, Modelling Assumptions, and Limitations	76
38.1 Primary Research Sources	76
38.2 Secondary Sources	77
38.3 Limitations	77
39 Appendices	77

1. Executive Summary

Global Capability Centers—the offshore delivery arms of the world’s leading multinationals—were built on a single premise: skilled human labour in emerging markets could perform the same cognitive work as Western employees at a fraction of the cost. For two decades, that premise delivered. It powered cities, created a professional middle class in countries that had lacked one, and allowed multinational corporations to extract margin from labour cost differentials that geography made possible and telecommunications made practical.

That premise is now being taken apart by artificial intelligence. The speed of the dismantling is faster, and the structural completeness of it more total, than most people working inside or around GCCs are prepared to accept.

Orbys360 spent twelve months surveying multiple global enterprises and tracking workforce composition in depth across 200+ GCCs across the globe. The finding is straightforward: more than 50% of current GCC headcount occupies roles with high-to-critical AI displacement risk within a five-year window.

“The competitive advantage of the GCC model—cheaper human cognition—is precisely what AI automates first, best, and at unlimited scale.”

—Orbys360, 2026

💡 ORBYS360 KEY FINDING

The greatest risk is not a sudden cliff. It is a slow-motion hollowing. GCC leaders report hiring freezes disguised as restructuring. Role counts are flat, but output expectations have doubled. The invisible workforce reduction is already underway. By the time it becomes undeniable in headlines, most organisations will be 18 months past the point of proactive response.

This report does not argue that AI adoption should be stopped. It argues that the transition, if managed with transparency and investment, can be navigated with far less human damage than a transition managed through silence, denial, or delay.

The 2.9 million professionals directly exposed to displacement are not an abstraction. They are individuals who made real financial and family decisions on the strength of career trajectories that may no longer exist. The obligation of the companies that employ them, the leaders who manage them, and the governments that depend on them is to meet that reality with honesty.

2. The GCC Model: A Twenty-Year Retrospective

The modern Global Capability Center emerged in the early 2000s as an evolution of the earlier BPO wave. Where BPO vendors were third-party service providers operating under contractual relationships, GCCs were captive operations—wholly owned subsidiaries of the parent company, staffed by its own employees, operating under its own governance frameworks.

The economics were simple. A software engineer in Bengaluru cost a quarter of what the same engineer cost in San Jose. An accountant in Warsaw cost a third of the London equivalent. A customer service representative in Manila cost a fifth of the rate in New York. The quality gap, which had been a legitimate concern in the early outsourcing era, narrowed rapidly as GCC operators invested in training, infrastructure, and management processes.

By 2010, GCCs had moved well beyond data entry and call centre work. They were performing financial modelling, writing production code, managing supply chains, conducting legal research, and building analytics capabilities that rivalled headquarters. The industry grew from a few hundred entities in 2005 to more than 1,700 in India alone by 2024, employing close to 1.9 million professionals in that country.

The economic model worked because it relied on a durable gap: the difference in the cost of human cognitive labour between geographies. Currency differentials, cost-of-living differences, and educational infrastructure in countries like India, Poland, and the Philippines created a steady supply of qualified professionals willing to work for wages that, while competitive locally, were far below Western equivalents.

What the industry did not appreciate—what it could not have appreciated until very recently—is that the same qualities that made GCC work attractive to parent companies also made it attractive to AI systems. The work was documented. It was standardised. It followed rules. It produced measurable outputs. It was, in a word, legible. And legibility is the precondition for automation.

“We spent twenty years turning ambiguous headquarters problems into clean, documented, repeatable processes. We wrote the playbooks. We standardised the outputs. We reduced variance. We were proud of it. We did not realise we were writing training manuals for the machines that would replace us.”

—GCC Managing Director, Financial Services

3. The Numbers: Aggregate Displacement Data

The following metrics represent Orbys360’s analysis across multiple global enterprises, with detailed workforce tracking of 200+ GCCs.

50%	2.9M	CRITICAL
GCC Roles Face Displacement	Professionals Directly Exposed	BPO/KPO Automation Risk Tier
High-to-critical AI risk within 2025–2031	Across surveyed GCC geographies	Highest tier of AI automation susceptibility
\$46B	HIGH	<2 YRS
Annual Labour Cost Under AI Threat	IT Development Roles Risk Tier	Average AI Deployment Lag
Replaceable over 3–5 year cycle	Code gen, testing, docs, Tier 1–2 support	From HQ approval to GCC reductions

The 50% figure is not a prediction of job losses within the 2027–2031 window. It is an assessment of structural exposure—the proportion of roles whose core task composition overlaps substantially with current or near-term AI capabilities.

The 2.9 million count does not include secondary employment effects in adjacent industries. In cities like Bengaluru and Manila, where GCC employment anchors entire local economies, the knock-on effects in housing, retail, transport, and food services could double or triple the affected population.

The revenue at stake is approximately \$110 billion—the estimated total annual revenue of the global GCC industry. The displacement risk does not threaten the entire revenue base, but it does threaten the cost structure and margin architecture on which the industry’s economics depend.

4. The Displacement Already Underway

The displacement projections in this report are not speculative. They describe the acceleration of a pattern that is already documented, named, and dated. Before examining the functional threat matrix or the technology driving displacement, it is worth establishing what has already occurred—because the transition from “could happen” to “is happening” took place before most GCC professionals were informed.

4.1 The Corporate Record

Accenture — 19,000 Jobs Cut (March 2023). Accenture disclosed in an SEC filing on 23 March 2023 that it would eliminate approximately 19,000 positions—2.5% of its 738,000-person workforce—at a restructuring cost of \$1.5 billion. The company did not officially attribute the cuts to AI. It simultaneously announced a \$3 billion investment in AI capabilities. Revenue grew approximately 4% year-on-year to \$64.1 billion in FY2023. By mid-FY2024, Accenture reported over \$3 billion in generative AI bookings. The workforce was recomposed: AI-related roles grew while traditional services headcount shrank.

IBM — 7,800 Back-Office Hiring Pause (May 2023). CEO Arvind Krishna told Bloomberg on 1 May 2023: “I could easily see 30% of that getting replaced by AI and automation over a five-year period.” The pause applied to HR, finance, and non-customer-facing functions. IBM’s total headcount declined from approximately 288,000 at end-2022 to 270,000–280,000 by late 2024. The reduction was implemented through attrition—positions vacated were simply not refilled.

TCS — Net Reduction of ~25,000 (Mid-2023 to Late 2024). Tata Consultancy Services peaked at approximately 616,000 employees in Q1 FY24 (June 2023) and declined to approximately 591,000 by Q3 FY25 (December 2024). Revenue grew 6–7% year-on-year throughout the period.

Infosys — Net Reduction of ~27,000 (FY2024). Headcount fell from approximately 343,000 to 316,000 over FY2024. Revenue was \$18.56 billion, up 1.4% year-on-year. FY25 growth guidance was raised multiple times.

Wipro — Net Reduction of ~29,000 (FY2024–Q2 FY2025). Headcount declined from approximately 259,000 to 230,000. Revenue was roughly flat, but margins improved.

Cognizant — Targeted Voluntary Separations (Late 2023). CEO Ravi Kumar referenced “pyramid optimisation” in Q3 2023 earnings. Over 3,500 separations occurred in a single round, targeting process and BPO roles.

Tech Mahindra — ~6,000 Headcount Reduction (FY2024). CEO Mohit Joshi publicly discussed “AI-led right-sizing” as a strategic priority.

Klarna — AI Replacing ~700 Customer Service Agents (2024). CEO Sebastian Siemiatkowski stated that AI was performing work equivalent to 700 customer service agents. Contractor and BPO staff reductions followed, heavily impacting offshore operations.

Dropbox — ~500 Roles Cut, 16% of Staff (2024). CEO Drew Houston explicitly cited AI as the reason for the reduction, affecting global and offshore teams.

Duolingo — 10% of Contractors Eliminated (January 2025). The company specifically cited AI replacing translation and content work previously performed by offshore contractors.

Microsoft — ~15,000 Roles Eliminated (2025). Microsoft cut approximately 15,000 positions across multiple rounds in 2025, with the largest single round in July eliminating around 9,000 roles. The company simultaneously committed tens of billions in AI infrastructure

spending. The pattern—substantial cuts in traditional functions while AI capacity expanded—mirrored the Accenture playbook from two years earlier.

Meta — ~11,600 Roles Cut Across Two Rounds (2025). Meta laid off approximately 3,600 employees in early 2025 and announced an additional 8,000 cuts later in the year. Public commentary linked the reductions to the company’s accelerated AI infrastructure spending and shift toward AI-augmented engineering teams.

TCS — 12,000 Layoffs Announced (July 2025). Tata Consultancy Services confirmed it would cut approximately 2% of its global workforce—over 12,000 employees—across FY25–26. CEO K. Krithivasan attributed the decision to a skills mismatch and limited feasibility of redeployment as AI-driven productivity gains compressed traditional service-line headcount needs.

Salesforce — 4,000 Customer Service Roles Cut (September 2025). CEO Marc Benioff stated publicly that the company had reduced its support workforce from 9,000 to 5,000 “because I need less heads,” attributing the decision directly to AI agents handling 50% of customer interactions and reducing support costs by 17%. Salesforce simultaneously moved to hire 3,000–5,000 sales roles, recomposing rather than shrinking total headcount.

Amazon — 14,000 Corporate Roles Eliminated (October 2025). Amazon confirmed plans to cut approximately 14,000 corporate employees as part of organisational changes the company described as AI-driven simplification. A further reduction of around 370 roles followed at its European headquarters in December 2025.

IBM — ~8,000 HR and Administrative Roles Eliminated (2025). IBM cut around 8,000 jobs in HR, finance, and other back-office functions over 2025, citing the deployment of AI tools handling routine administrative tasks—a delayed but explicit follow-through on CEO Arvind Krishna’s 2023 statement that 30% of such roles could be replaced by AI within five years.

Cognizant — Net Reduction of ~13,700 (FY25 Trailing). Cognizant headcount declined from approximately 350,000 in Q1 FY25 to 336,300 over the course of the year, the steepest contraction among major IT services firms tracked. Revenue continued to grow.

4.2 The Industry Pattern

💡 ORBYS360 KEY FINDING

Across the top five Indian IT services firms alone, the combined net reduction was approximately 81,000 employees between mid-2023 and late 2024—while aggregate revenue held steady or grew. NASSCOM data for FY2024 confirmed the first aggregate hiring decline across top-tier Indian IT firms in over a decade. The pattern accelerated in 2025: a further ~25,000 net reduction across Tier-1 IT firms, more than 50,000 cuts at the largest US tech employers, and explicit AI attribution in CEO statements that had been absent from earlier rounds. Revenue went up. Headcount went down. That differential is the displacement already happening.

The BPO sector is experiencing similar pressure from the client side. Teleperformance’s share price dropped approximately 30% in 2023 on AI-related concerns; the company acknowledged that AI was reducing call volumes and headcount requirements across accounts. In 2025, both Teleperformance and Concentrix announced AI-first strategic plans (Teleperformance’s “Future Forward” programme, Concentrix’s Solv platform) explicitly designed to deliver client outcomes with reduced human labour intensity. TaskUs warned investors in 2024 that generative AI adoption by clients constituted a material business risk; that risk crystallised across BPO P&Ls in 2025.

Everest Group estimated in 2024 that AI could automate 20–30% of offshore IT services work by 2027. Gartner projected that AI would reduce demand for offshore IT services roles by 15–20% by 2026.

“None of these companies announced ‘AI layoffs.’ They announced restructuring, optimisation, and efficiency programmes. The AI was the cause. The language was the cover. The pattern—revenue up, headcount down—is the evidence.”
—Orbys360, 2026

5. The Real Economics: Cost Arbitrage Made Visceral

The \$46 billion at risk is an aggregate. Aggregates do not change CFO behaviour. Unit economics do.

The GCC model was built on a simple cost equation: a skilled professional in an emerging market costs 25–33% of the Western equivalent. That equation was compelling for two decades. The AI equation is more compelling by an order of magnitude.

5.1 The Math That Changes Boardrooms

Role	Annual Cost	AI Alternative	AI Coverage	Payback
Junior Developer (Bengaluru)	\$18,000/yr	GitHub Copilot Enterprise \$390/seat/yr	Majority of output	Weeks
BPO Customer Service Agent (Manila)	\$8,000/yr	AI Voice Agent (GPT-4o class)	Most queries	~2 weeks
Financial Analyst (Warsaw)	\$28,000/yr	LLM analytics agent	Majority of analysis	~1 month
QA Tester (Hyderabad)	\$14,000/yr	AI testing suite	Most test cases	Weeks
Data Entry Clerk (Cairo)	\$6,000/yr	Document AI (IDP)	Near-total volume	Days
HR Operations Analyst (Pune)	\$12,000/yr	AI payroll + HR chat-bot	Vast majority of queries	~2 weeks

Note: AI coverage indicates how much of the role's task output AI systems can currently handle at comparable quality. Payback period = time for AI cost savings to exceed annual AI tooling cost.

AI coverage is not total. AI does not do the whole job. But a team of twenty contracts to a fraction of its prior size. Those who remain are more skilled, better paid, and more productive—but the eliminated positions do not return. Multiply this across 1,700 GCCs and the structural arithmetic becomes clear.

5.2 The Team-Level View

A GCC finance operations team of 25, performing AP/AR processing, reconciliation, and management reporting at a fully-loaded cost of ~\$450,000/year, can be reduced to 8 professionals managing AI-driven workflows at a team cost of ~\$200,000/year plus ~\$15,000/year in AI tooling. The parent company saves \$235,000/year on a single team. Across 200 such teams in a large financial services GCC, the arithmetic produces \$47 million in annual savings. The payback on AI deployment: under 90 days.

The unit economics explain why the displacement pattern documented in the previous section is accelerating rather than stabilising. Every CFO who sees these numbers reaches the same conclusion. The only variable is timing.

ORBYS360 STRATEGIC INSIGHT

The GCC cost advantage was 3–4× over Western equivalents. The AI cost advantage is 20–50× over GCC equivalents. The arbitrage that built the GCC industry has been superseded by a larger arbitrage that threatens to dismantle it. The question is not whether the economics justify the transition. The question is how fast the transition will proceed and who will manage it responsibly.

6. Functional Threat Matrix

Orbys360’s activity-level analysis decomposes GCC work into three primary functional domains—Business Process Management (BPM), IT Activities, and Engineering Research & Development (ER&D). Each domain is further broken down into clusters, sub-clusters, and individual activities, with risk level, displacement horizon, AI mechanism, and residual human value assessed for every activity. The full activity-level breakdown is presented in Appendix E; this section presents the cluster-level summary and key findings.

Domain	Risk	Horizon	Key AI Mechanisms	Residual Value	Human
BPM	CRITICAL	2025–28	Invoice AI, AI voice agents, document AI, RegTech, LLM contract generation, HR chatbots	Exception judgment, regulatory interpretation, complex negotiation, audit sign-off	
IT Activities	HIGH	2025–30	AI code gen (Copilot, Cursor), AIOps, ITSM bots, AI-SIEM/SOAR, LLM auto-docs	Architecture decisions, security design, complex systems, strategic analytics	
ER&D	MEDIUM	2025–30	Generative CAD, physics-informed ML, AI research agents, LLM SOP generation	Creative judgment, safety certification, novel failure modes, IP strategy	

6.1 BPM — Business Process Management

BPM activities represent the highest-risk functional cluster in GCCs. AI agents now perform data entry, invoice processing, customer query handling, compliance verification, and contract abstraction at superhuman speed. Most BPM sub-functions reach critical risk by 2026–27. Voice AI has eliminated the last defensible BPM moat: the human touch in customer interactions.

Sub-Function	Risk	Horizon	Primary AI Mechanisms	Residual Human Value
Finance & Ac-counting	CRITICAL	2025–28	Invoice AI (OCR+NLP) end-to-end, AI cash application, predictive collections, automated journal entries, AI-driven reconciliation, RegTech auto-filing, AI cash forecasting, automated workpapers	Vendor escalation, judgmental accruals, audit sign-off, novel tax positions, counterparty relationships, cross-border tax strategy
HR Operations	CRITICAL	2025–29	End-to-end AI payroll engines, AI resume screening & chatbot scheduling, AI-driven HRIS updates, AI eligibility verification, HR chatbots resolving Tier 1–2 queries, automated offboarding workflows	Legal exceptions, cultural fit in senior hiring, sensitivity handling, complex life events, manager coaching, complex grievances
Customer Ops	CRITICAL	2025–28	AI voice agents (GPT-4o/Gemini Live) handling end-to-end calls, LLM-powered chat with CRM integration, document AI for KYC, AI data enrichment & deduplication, AI collection agents with predictive dialler	VIP escalations, regulatory complaints, high-risk customer judgment, STR filing, data governance, hardship assessment
Procurement & Supply Chain	CRITICAL	2025–29	Autonomous PO creation, end-to-end invoice AI with automated 3-way matching, AI vendor validation & auto-onboarding, contract AI for milestone tracking, AI spend categorisation & anomaly detection, AI demand forecasting	Strategic supplier decisions, fraud exceptions, vendor relationship governance, renegotiation, strategic sourcing, crisis logistics
Compliance & Risk	CRITICAL	2025–28	Document AI for KYC/AML, automated real-time sanctions & PEP screening, AI false-positive reduction, RegTech auto-reporting, AI-generated audit logs, automated DSAR processing, AI control testing scripts	Complex risk decisions, STRs, novel regulatory interpretations, external auditor liaison, novel privacy interpretations, risk appetite decisions
Legal Operations	CRITICAL	2025–27	LLM contract generation from playbooks, document AI for clause extraction & obligation mapping, LLM legal research (case law, precedents), AI contract review with redline generation, AI patent search & prior art mining	Novel commercial terms, ambiguous multi-jurisdiction contracts, strategic legal advice, unusual risk positions, IP strategy
Data Entry & Document Mgmt	CRITICAL	2025–27	End-to-end document AI (zero human-in-the-loop for standard docs), multimodal AI OCR, AI auto-classification & meta-data tagging, Forms AI (Intelligent Document Processing), LLM auto-generation of SOPs	Disputed/handwritten edge cases, degraded/rare-language documents, legal hold decisions, sensitive/regulatory tone review

A single AI agent can process invoices at a rate equivalent to 40–60 human operators, with error rates that are, in many documented deployments, lower than the human baseline. The last defensible position for human BPO operators was voice-based customer service. Voice AI

has eroded that position rapidly. In 2024, several major enterprises reported that customers could not reliably distinguish between AI voice agents and human agents in blind testing.

The World Economic Forum projects that GCC CFO function headcount will decline by 40% by 2029. Orbys360's modelling shows the steepest near-term displacement in reconciliation and compliance roles, followed by management reporting and financial planning support. Across BPM, the residual human value concentrates in exception handling, regulatory judgment, and relationship-dependent decisions—work that requires accountability, empathy, or the ability to navigate ambiguous regulatory terrain.

6.2 IT Activities

IT activities face a bifurcated threat landscape. IT Support (L1) and QA roles face critical displacement by 2025–26 as AI agents fully automate routine helpdesk and testing tasks. Application development is under sustained pressure from AI coding agents, with junior-to-mid developer roles at high risk by 2026–29. Infrastructure management faces a slower but structural shift to AIOps. Senior architects, security specialists, and complex systems engineers retain the highest residual value.

IT Cluster	Risk	Horizon	Primary AI Mechanisms	Residual Human Value
Application Development	HIGH	2025–30	LLM user story generation, AI code gen (Copilot, Cursor) for CRUD/API/boilerplate, AI design tools (Figma AI, Galileo), AI visual regression & exploratory testing, AI test script auto-generation, AI release orchestration, LLM auto-docs (Mintlify, Swimm)	Stakeholder strategy, domain ambiguity resolution, creative direction, architecture decisions, complex business logic, security design, test strategy, multi-cloud governance
Infrastructure Mgmt	HIGH	2026–31	AIOps automated alerting with self-healing scripts, AI IaC generation (Terraform), AIOps platforms (Dynatrace, Datadog AI) with autonomous remediation, AI query optimisation, AI-SIEM with autonomous alert triage, SOAR playbook automation, AI discovery & auto-inventory	Novel outage patterns, complex multi-vendor architectures, architecture design, cost strategy, schema design, advanced persistent threat response, licensing negotiation
IT Support & Service Desk	CRITICAL	2025–30	Full ITSM bot automation (zero human touch for standard requests), AI chat agents with system integration, AI auto-classification & priority routing, AI diagnostic agents & auto-runbook execution, AI incident orchestration, LLM-powered learning assistants	Complex identity edge cases, unusual software requests, ambiguous escalations, complex cross-app integration issues, war-room leadership, change management
Data & Analytics Eng	HIGH	2026–29	AI code gen for data pipelines, AI data quality profiling & anomaly detection, AI auto-metadata tagging & lineage generation, AI-generated dashboards, natural language BI query engines, LLM data analysis agents	Complex data architecture, business data governance policy, enterprise data strategy, strategic analytics narrative, hypothesis framing

Software development displacement is concentrated in the junior-to-mid tier because it includes significant components that remain resistant to automation. Architecture decisions, system design, debugging complex production issues, and translating ambiguous business requirements into technical specifications all require contextual judgment that current AI systems handle poorly. Senior engineers are more productive with AI tools, not less needed. The split runs through the middle of most GCC development teams.

The IT Support cluster faces the most immediate displacement. Helpdesk agents handling password resets, software installation, and FAQ queries are fully automatable today—AI chat agents with system integration handle ticket-to-resolution without human intervention. Tier 2 application support retains more value, but AI knowledge retrieval and diagnostic agents are rapidly compressing the need for human intermediaries.

6.3 ER&D — Engineering Research & Development

ER&D activities face the most stratified threat. Routine and documentation-heavy roles—technical manuals, test reports, literature reviews, CAD drafting—are approaching critical displacement by 2025–27. Simulation and analysis roles face high-to-medium risk as physics-informed ML models replace standard solver runs. Senior design engineers, research scientists, and certification specialists retain high residual value—their creative judgment, regulatory accountability, and domain expertise remain outside current AI capability. The displacement horizon for ER&D is slightly longer than BPM, but the trajectory is structurally inevitable.

ER&D Cluster	Risk	Horizon	Primary AI Mechanisms	Residual Human Value
Product Design & Engineering	HIGH	2026–29	AI literature mining, generative AI DFM analysis, AI impact analysis for ECOs, generative CAD AI (auto-generates 2D from 3D), AI-assisted parametric modelling, AI PLM data cleansing & BOM validation, AI standards mapping against design parameters	Creative direction, supplier partnership, cross-functional negotiation, customer specification interpretation, complex assembly design, configuration governance, authority liaison
Simulation & Analysis	MEDIUM	2026–30	AI-driven FEA setup & mesh automation, AI load case generation, AI-augmented topology optimisation, physics-informed ML replacing full CFD for standard cases, AI surrogate models for thermal prediction, AI EMC simulation, AI-assisted control system modelling	Failure mode judgment, material uncertainty, manufacturing feasibility, novel flow regimes, complex multi-physics, test lab coordination, safety certification
Testing & Validation	HIGH	2025–29	AI test planning from requirements, AI sensor data processing & anomaly detection, LLM auto-generation of test reports from raw data, AI Weibull analysis automation, AI-assisted digital twin calibration, AI FMEA generation from design data, AI regulatory cross-reference checking	Risk prioritisation, experimental judgment, certification sign-off, warranty strategy, model validation, novel failure modes, authority relationships
Research & Development	HIGH	2026–30	AI research agents (Semantic Scholar, Perplexity Deep Research), AI patent landscape analysis, AI rapid prototyping assistants, AI statistical modelling, LLM grant proposal drafting, AI schedule & resource optimisation, AI automated status reports	Research hypothesis framing, strategic investment decisions, IP strategy, creative problem framing, intellectual originality, stakeholder alignment, board-level decisions
Tech Docs & Knowledge Mgmt	CRITICAL	2025–28	LLM end-to-end SOP generation, AI IETP & illustrated parts catalogue automation, AI eLearning module creation, AI document automation from design data, AI auto-curating & deduplicating knowledge articles	Safety-critical tone, OEM review & liability sign-off, instructional design strategy, contractual accuracy, knowledge architecture strategy, cross-programme synthesis

Engineering functions occupy a structurally different position from BPM and IT. The work frequently involves interaction with physical systems, regulatory certification requirements, and consequences of failure—structural, safety, environmental—that impose a higher bar for AI adoption. The exposed portion consists primarily of routine engineering support: preparing CAD drawings from specifications, setting up simulation runs, processing test data, writing standard engineering reports, and performing routine design validation. The resistant portion includes core engineering judgment: selecting materials for novel applications, analysing failure modes in complex systems, and making design decisions where the consequences of error are severe.

The analytics and KPO functions—previously a distinct GCC category—are distributed across BPM (compliance analytics, financial reporting) and IT (data engineering, BI reporting). The work product that a team of five research analysts might produce over a week—a competitive landscape report, a market sizing exercise, a regulatory environment scan—can now be generated by an AI system in hours, at a quality level comparable for routine analyses. The proportion of KPO work that falls into the “routine” category is larger than most KPO professionals would like to admit.

7. Role-Level Displacement Analysis

The following table presents Orbys360’s role-level displacement assessment, drawn from detailed tracking of 200+ GCCs across the global survey universe, mapped against AI mechanisms and residual human value. Methodology is detailed in Section 38.

Role Category	Horizon	Risk	Primary AI Mechanism	Residual Human Value
Data entry & processing clerks	2025–26	CRITICAL	End-to-end document AI; zero human-in-the-loop for standard documents. Multimodal OCR, Forms AI (IDP).	Disputed/handwritten edge cases, degraded or rare-language documents
Customer service reps	2025–27	CRITICAL	AI voice agents (GPT-4o/Gemini Live) handling end-to-end calls, LLM-powered chat with CRM integration	Emotionally complex interactions, VIP escalations, regulatory complaints
Finance & accounts processing	2025–27	CRITICAL	Invoice AI (OCR+NLP), automated reconciliation, AI-driven journal entries, RegTech auto-filing	Judgmental accruals, audit sign-off, novel tax positions, cross-border strategy
HR operations & payroll	2025–27	CRITICAL	End-to-end AI payroll engines, AI resume screening, HR chatbots resolving Tier 1–2 queries	Legal exceptions, sensitive conversations, complex grievances, cultural fit
Technical documentation	2025–26	CRITICAL	LLM auto-generation from code, API specs, engineering inputs (Mintlify, Swimm)	Safety-critical tone, regulatory approval, external-facing narrative
Junior software developers	2026–28	HIGH	AI code gen (Copilot, Cursor) for CRUD, API scaffolding, boilerplate, UI components	Complex business logic, security design, architecture decisions
QA & testing engineers	2026–27	HIGH	AI visual regression, AI-driven exploratory testing agents, AI auto-generating test scripts	High-risk edge case validation, test strategy, coverage architecture
Research analysts & KPO	2026–28	HIGH	LLM data analysis agents, auto-insight generation, AI patent landscape analysis	Hypothesis framing, business context, strategic investment decisions
Procurement & supply chain	2025–29	HIGH	Autonomous PO creation, end-to-end invoice AI with 3-way matching, AI demand forecasting	Strategic supplier decisions, crisis logistics, fraud exceptions
Engineering support & CAD	2026–29	MODERATE	Generative CAD AI (2D from 3D), AI-assisted parametric modelling, AI FEA setup	Complex assembly design, customer spec interpretation, failure mode judgment
Mid-level IT infra ops	2026–29	MODERATE	AIOPs (Dynatrace, Datadog AI) with autonomous remediation, AI IaC gen (Terraform)	Platform architecture, cost strategy, vendor negotiation

The roles with the highest displacement risk—data entry, customer service, and technical documentation—also have the shortest time horizons. These are 2025–2026 problems, not 2030 problems. Some of the displacement is already happening, visible in the hiring freezes and unreplaced attrition documented in Section 4 rather than in formal layoff announcements.

Junior software developers represent the single largest role category by headcount share. Their high displacement risk over a 2026–2028 horizon means that the largest cohort of GCC professionals is facing structurally severe pressure. This group is also the pipeline from which senior developers, architects, and engineering leaders are drawn. If the pipeline compresses, the downstream effects on GCC capability will be felt for years.

8. The Technology Driving Displacement

Understanding the specific AI capabilities that are driving GCC displacement is necessary for assessing which roles are most exposed and on what timeline. The AI systems relevant to GCC operations fall into four categories, each at a different stage of enterprise readiness.

Large Language Models and Cognitive Task Automation. Large language models—the technology behind ChatGPT, Claude, Gemini, and their enterprise equivalents—reached human parity on knowledge work benchmarks in 2022–2023. They can summarise documents, draft reports, answer questions about regulatory frameworks, write code, and produce analysis that is, for routine tasks, indistinguishable from the output of a mid-level professional.

For GCCs, the significance of LLMs is that they commoditise the exact type of work that GCCs were built to deliver: structured cognitive output produced by trained professionals following established processes. When a financial analyst in a GCC spends three days preparing a quarterly management report, and an LLM can produce a comparable report in minutes using the same source data, the economic calculus changes.

AI Coding Agents. AI coding tools—GitHub Copilot, Amazon CodeWhisperer, and newer agentic coding systems—have moved beyond autocomplete to autonomous code generation. In 2024 benchmarks, leading AI coding systems could generate production-quality code for well-defined tasks (API endpoints, database queries, front-end components, unit tests) with success rates exceeding 70%.

The junior-to-mid-level developer who writes boilerplate code, creates test suites, and documents APIs is performing work that AI coding agents can now handle. The developer who architects systems, debugs complex failures, and translates ambiguous business requirements into technical designs is performing work that AI cannot yet handle reliably. The split runs through the middle of most GCC development teams.

Robotic Process Automation and Intelligent Document Processing. Traditional RPA was brittle—it followed hard-coded rules and broke when interfaces changed. The integration of AI with RPA has produced a new generation of intelligent automation that can handle unstructured data, adapt to interface changes, and make judgment calls within defined parameters.

Invoice processing that required a human to interpret handwritten notes or resolve discrepancies can now be handled by AI-enhanced RPA systems with accuracy rates above 95%.

Agentic AI Systems. The most consequential development is the emergence of agentic AI—systems that operate autonomously across multi-step workflows, making decisions, executing actions, and handling exceptions without human intervention. Before agentic AI, an AI system could help a human do a task faster. With agentic AI, the system owns the task end-to-end. The human, if involved at all, is reviewing the output rather than doing the work.

Enterprise adoption of agentic AI began in earnest in 2025 and accelerated in early 2026. The systems are maturing rapidly, and reliability is improving with each deployment cycle. The GCC functions most exposed—BPO, financial processing, and routine analytics—are the same functions that make up the largest share of GCC employment.

9. Sector Analysis: Financial Services GCCs

Financial services GCCs are the largest single sector in the global GCC industry by revenue and among the most exposed to AI displacement. The major banks—JPMorgan Chase, Goldman Sachs, Citigroup, HSBC, Barclays—each operate GCCs with headcounts ranging from 10,000 to over 50,000 employees.

The functions performed by financial services GCCs map directly onto the highest-risk categories in the BPM threat matrix. Accounts payable processing is being automated end-to-end by invoice AI systems combining OCR and NLP agents. Accounts receivable is compressed by AI cash application and predictive collections with auto-dunning. Reconciliation and month-end close are handled by automated journal entries and AI-driven reconciliation engines. Financial reporting is produced by auto-report generation with narrative AI. Tax compliance is processed by RegTech AI with auto-filing and tax rule engines. Audit support is accelerated by AI evidence gathering and automated workpapers.

The residual human value in these functions concentrates in vendor escalation and exception judgment (AP), complex dispute resolution (AR), judgmental accruals and audit sign-off (GL), strategic commentary and board narrative (reporting), novel tax positions and jurisdiction nuance (tax), and auditor judgment and risk assessment (audit).

Compliance and risk functions face equally severe pressure. Document AI handles KYC/AML processing. Automated real-time sanctions and PEP screening with AI false-positive reduction compresses what were large teams. RegTech AI automates regulatory reporting and submission. AI-generated audit logs and automated evidence compilation reduce audit trail teams. The residual value lies in complex risk decisions, suspicious transaction reports, novel regulatory interpretations, and regulatory dialogue—work requiring regulatory judgment and accountability.

Internal pilot programmes at several major banks have documented results that confirm the modelling. At one large US bank, an AI system was tested against a GCC team performing loan document review. The AI system processed documents at 12 times the speed of the human

team, with a 3% lower error rate. The pilot was classified internally and the GCC leadership was not informed of the results.

The pattern of classified pilots deserves attention. The information asymmetry it creates leaves GCC employees unable to make informed career decisions.

Near-term trajectory. Orbys360 projects that financial services GCCs will experience the earliest visible headcount reductions, beginning in 2026. The combination of high displacement probability, strong HQ cost scrutiny, and relatively mature AI tooling for financial functions means that the conditions for contraction are already in place. The first reductions will appear as continued non-backfilling of attrition, followed by formal restructuring announcements from 2027 onward.

Insurance GCCs face a particularly compressed timeline. Claims processing—the bread and butter of insurance GCC operations—is among the most automatable functions in the financial services sector. AI claims systems are already in production at several major insurers, and the performance data is unambiguous. The question for insurance GCC leaders is not whether to prepare for displacement, but whether they still have time to do so.

10. Sector Analysis: Tech GCCs

Technology company GCCs present a more complex displacement picture than financial services. The work is more varied, the workforce skews younger and more technically skilled, and the parent companies are themselves the developers of AI technology.

The displacement risk maps directly onto the IT Activities threat matrix. In **Application Development**, AI code generation tools (Copilot, Cursor) write standard UI components, CRUD operations, API scaffolding, and boilerplate at the junior level. AI design tools (Figma AI, Galileo) generate wireframes and prototypes from prompts. AI auto-generates test scripts from user flows and performs visual regression and exploratory testing autonomously. LLM tools (Mintlify, Swimm) auto-generate technical and API documentation from code and specs. The residual human value lies in architecture decisions, complex business logic, security design, creative direction, and test strategy.

In **Infrastructure Management**, AIOps platforms (Dynatrace, Datadog AI) provide autonomous remediation, while AI-SIEM with autonomous alert triage and SOAR playbook automation compress SOC L1 teams. AI IaC generation (Terraform) and AI-generated Helm charts reduce junior cloud and container operations roles. AI query optimisation and automated performance tuning reduce DBA requirements. The residual value lies in novel outage patterns, advanced persistent threat response, platform architecture, and complex multi-vendor architectures.

In **IT Support**, ITSM bots achieve full automation for standard requests—password resets, software installation, FAQ resolution—with zero human touch. AI auto-classification, priority assignment, and routing handle ticket management. AI diagnostic agents with auto-runbook

execution compress Tier 2 application support. The residual value lies in complex identity edge cases, unusual software requests, cross-app integration issues, and war-room leadership during major incidents.

The less exposed areas remain senior engineering, platform architecture, security engineering, and product management. These roles require systemic thinking, judgment under ambiguity, and cross-functional coordination that AI systems do not yet replicate.

An important dynamic in technology GCCs is that the parent companies are often enthusiastic AI adopters across their own operations. The GCC is building AI products for customers while simultaneously being displaced by AI in its own internal operations. Several technology company GCC leaders described this to Orbys360 as an uncomfortable dissonance that their organisations have not yet resolved.

A QA team of twenty that manually tests software releases can, with AI testing tools, be reduced to a team of five or six managing and reviewing automated test runs. The humans who remain are more skilled and higher-paid than the team they replaced, but there are far fewer of them. This pattern—fewer, better, more expensive humans supervising AI systems—is the structural template for what GCC technology delivery will look like by 2029.

The talent pipeline effects are worth noting. Technology GCCs have traditionally served as training grounds for junior engineers, who spend two to three years building foundational skills before progressing into senior roles. If the junior entry point compresses, the pipeline that produces senior engineers, architects, and engineering leaders will narrow. The downstream effects on GCC capability will be felt for years, even by organisations that manage the transition well.

11. Sector Analysis: Healthcare and Life Sciences GCCs

Healthcare and life sciences GCCs face significant AI automation potential in back-office functions—claims processing, medical coding, revenue cycle management, pharmacovigilance reporting—but are subject to regulatory constraints that slow the pace of AI deployment relative to other sectors.

Medical coding, a major function in healthcare GCCs, faces critical displacement risk. AI medical coding systems can assign ICD-10 codes to clinical documentation with accuracy matching or exceeding human coders. But CMS guidelines in the US and MDR in the EU require human oversight of AI-generated coding decisions, limiting the pace of headcount reduction. The residual human value lies in regulatory sign-off and novel coding scenarios—precisely the kind of accountability-dependent work that slows but does not prevent automation.

Pharmacovigilance is another high-exposure function. Processing adverse event reports—reading clinical narratives, classifying events by severity, entering data into regulatory databases—is work that AI systems now perform with high accuracy. Document AI handles clinical document processing, while Forms AI (Intelligent Document Processing) automates medical forms

across all types. Several pharmaceutical companies have deployed AI pharmacovigilance tools in their GCC operations, reducing per-report processing time from 30–45 minutes to under 5 minutes. Compliance operations face pressure from automated DSAR processing and AI consent tracking. The residual value remains in novel privacy interpretations and complex regulatory adjudication.

On the R&D side, AI research agents (Semantic Scholar API, Perplexity Deep Research) are compressing literature review and state-of-the-art survey work. AI statistical modelling automates insight extraction from clinical trial datasets. LLM auto-generation of test reports from raw data and templates handles standard regulatory documentation. The residual human value in healthcare ER&D lies in research hypothesis framing, certification sign-off, and safety integrity judgment.

Clinical trial data management, regulatory writing, and medical affairs support are moderately exposed. The regulatory and safety implications of errors impose a higher bar for AI adoption.

The net effect for healthcare GCCs is a slower but still substantial displacement trajectory. Orbys360 estimates that healthcare GCCs will experience 18–24 months of regulatory delay relative to financial services GCCs before displacement becomes visible in headcount numbers. That delay is a window for proactive transition, and healthcare GCC leaders who use it well will be better positioned than their financial services counterparts, who face a more compressed timeline.

12. Sector Analysis: Manufacturing GCCs

Manufacturing GCCs are concentrated in supply chain management, procurement, engineering support, and quality management. The global manufacturing sector operates more than 400 GCC entities, with the largest concentrations in India (Bangalore, Pune, Chennai) and Poland (Wroclaw, Lodz).

Supply chain and procurement functions face high displacement risk, driven by AI systems mapped in the BPM threat matrix. Autonomous PO creation with supplier matching AI and ERP integration bots handles routine purchase order processing. End-to-end invoice AI with automated PO/GR/Invoice 3-way matching processes the highest-volume procurement activity. AI vendor data validation, deduplication, and auto-onboarding compresses vendor master management. Contract AI handles milestone tracking, auto-renewal alerts, and clause extraction. AI spend categorisation with auto-dashboard generation and anomaly detection replaces routine analytics. AI demand forecasting models with automated replenishment signals reduce planning teams. The residual human value concentrates in strategic supplier decisions, fraud exception handling, vendor relationship governance, renegotiation and legal interpretation, strategic sourcing recommendations, and crisis logistics.

On the engineering side, the ER&D threat matrix reveals a stratified picture. Generative CAD AI auto-generates 2D drawings from 3D models, compressing the drafting function. AI-assisted parametric modelling and geometry generation from sketches reduces junior-to-mid 3D CAD

roles. AI PLM data cleansing, BOM validation, and automated release workflows compress product lifecycle management. AI-driven FEA setup, mesh automation, and results auto-interpretation handle routine structural analysis. Physics-informed ML models are replacing full CFD solver runs for standard simulation cases. AI sensor data processing with automated anomaly detection handles test data acquisition. LLM auto-generation of test reports from raw data handles standard documentation. The residual value lies in customer specification interpretation, complex assembly design, failure mode judgment, novel flow regimes, and safety certification.

Engineering support and quality management retain the lowest overall risk because they require hands-on interaction with physical systems, regulatory certification requirements, and consequences of failure that impose a higher bar for AI adoption. But the data analysis, reporting, and documentation components of these roles—which can account for 30–40% of the working week—are automatable by the mechanisms described above.

Manufacturing GCCs also perform significant back-office work in finance, HR, and IT support that faces the same displacement pressures as those functions in any other sector. The manufacturing-specific displacement is incremental to the cross-sector displacement that affects all GCCs.

13. Sector Analysis: Retail GCCs

Retail GCCs perform merchandising analytics, demand forecasting, customer analytics, e-commerce operations, and increasingly, personalisation at scale. The global retail sector operates approximately 200 GCC entities, with the largest operations belonging to companies like Walmart, Target, Tesco, and major fashion retailers.

The analytical functions are highly exposed. AI can generate merchandising recommendations, forecast demand with greater accuracy than human analysts, segment customers at a granularity that manual analysis cannot match, and optimise pricing in real time. A team of twenty demand forecasting analysts in a retail GCC is performing work that AI systems can now do faster, at larger scale, and with measurably better accuracy on structured historical data.

Customer analytics—once the jewel in many retail GCC mandates—is undergoing a particular compression. AI-generated dashboards, natural language BI query engines, and LLM data analysis agents with auto-insight generation now produce the customer segmentation models, lifetime value calculations, and campaign attribution analyses that GCC analytics teams spent weeks creating. AI data enrichment and deduplication agents with CRM auto-update compress customer data management teams. The residual human value lies in hypothesis framing, business context interpretation, and translating analytical findings into actionable business recommendations—work that requires understanding of the retail business in ways that pure data analysis does not.

Where retail GCCs may find durable advantage is in localisation and market-specific intelligence. A retail GCC in India or Poland that understands local consumer behaviour, regulatory requirements, and competitive dynamics can provide strategic intelligence that is not easily replicated by AI systems trained on global datasets. The challenge is that this kind of work requires a fundamentally different mandate than the analytics-at-scale model on which most retail GCCs were built.

14. Sector Analysis: Energy GCCs

Energy GCCs are concentrated in financial operations, HSE (health, safety, and environment) data management, engineering support, and increasingly, sustainability reporting and carbon accounting. The sector includes GCC operations for major oil and gas companies, utilities, and renewable energy firms.

The financial operations within energy GCCs are exposed at rates similar to the financial services sector—accounts processing, reconciliation, and management reporting face critical displacement risk. These functions are not energy-specific; they are general financial operations performed in an energy context, and they face the same AI displacement as their counterparts in banking or insurance.

HSE data management is moderately exposed, constrained by the regulatory and safety implications of errors. AI-generated audit logs and automated evidence compilation handle routine HSE documentation. AI auto-classification with retention policy enforcement and metadata tagging compresses records management. Processing safety incident reports, tracking compliance metrics, managing environmental monitoring data—these are tasks that AI can perform with high accuracy, but where the regulatory framework requires human review and sign-off. The residual human value lies in sign-off authority, external auditor liaison, and legal hold decisions. The human role is shifting from data processing to quality assurance and regulatory interpretation.

Engineering support remains the least exposed function in the energy sector. Upstream oil and gas engineering—reservoir modelling, well design, production optimisation—requires specialised domain knowledge, physical intuition, and an understanding of subsurface conditions that are inherently uncertain. However, the ER&D threat matrix shows that even here, physics-informed ML models are replacing full CFD solver runs for standard flow simulation cases, AI surrogate models handle routine thermal and heat transfer analysis, and AI auto-report generation from simulation results compresses post-processing teams. The residual value lies in complex multi-physics analysis, safety-critical design decisions, and novel flow regimes. The consequences of errors in petroleum engineering can be catastrophic (blowouts, environmental disasters), which imposes a level of conservatism about AI adoption that other sectors do not share.

Renewable energy GCCs are a growing but distinct segment. Wind farm performance analytics, solar yield modelling, and grid integration analysis are tasks where AI is increasingly capable, but where the sector is growing fast enough that AI augmentation is creating efficiency gains

rather than displacing workers. The renewable energy GCC is, for now, one of the few segments where AI is more likely to accelerate hiring than to reduce it.

Across all three sectors—manufacturing, retail, and energy—a common pattern is visible: the back-office and analytical functions that GCCs were originally established to perform are the most exposed, while the domain-specific, judgment-intensive functions that some GCCs have grown into are the least exposed. The GCCs that diversified up the value chain over the past decade are structurally more resilient than those that remained focused on volume-based processing.

15. The Timeline: From Language Models to Agentic Systems

The GCC industry has weathered technology transitions before—the shift from mainframes to client-server, the rise of cloud computing, the adoption of DevOps. None of these eliminated the GCC model. The AI transition is different in kind.

Year	Phase	Description
2022	<i>The capability threshold</i>	Large language models achieved human parity on knowledge work benchmarks. Legal reasoning, financial summarisation, code generation—tasks that GCCs charged premium rates to deliver—became available through APIs for fractions of a cent per task.
2023	<i>Proof of concept</i>	Parent companies began running pilot programmes at scale. Orbys360’s survey of multiple global enterprises found internal reports documenting 40–80% task-equivalence at 2–5% of the cost. Pilot programmes were classified. GCC leadership was not informed.
2024	<i>The quiet freeze</i>	Hiring volumes at major GCCs dropped 22% year-on-year despite record parent company revenues. Attrition was not backfilled. AI infrastructure was silently replacing incremental headcount.
2025	<i>Agentic AI arrives</i>	AI agents—systems that operate autonomously across multi-step workflows—became enterprise-ready. For the first time, AI owned the task from input to output.
2026–31	<i>The reckoning</i>	Orbys360 projects active GCC headcount contractions to begin appearing in public disclosures from 2026 onward. First movers will be cost-centre GCCs in financial services.

“The GCC industry spent twenty years making itself legible to management—and in doing so, made itself legible to machines. The very rigour that drove margins will now drive displacement.”

—Orbys360, 2026

16. The Acceleration Curve: AI Capability vs. GCC Hiring Cycle

One structural fact is more alarming than any displacement percentage: AI capability improvement is now faster than the GCC employment lifecycle.

16.1 The Timing Mismatch

Metric	Duration
Average time to approve, recruit, hire, and onboard a GCC role	9–12 months
AI capability doubling time on benchmark tasks (coding, analysis, document processing)	6–9 months
Average GCC budget cycle (from business case to operational deployment)	12–18 months
Time from AI pilot approval to production deployment in enterprise settings	3–6 months

The implications are structural. On SWE-bench coding benchmarks, AI pass rates improved from approximately 14% in early 2024 to over 50% by early 2026. Legal document review accuracy improved from 78% to 94% in eighteen months. Invoice processing, customer query handling, and test automation have seen comparable acceleration curves.

CRITICAL ADVISORY

By the time a GCC completes a hiring cycle for a role, the AI that was 50% capable at the job posting date is 80%+ capable at the joining date. The workforce planning models that GCCs use assume a stable capability baseline for the roles they are filling. That assumption is now false. Every hiring decision made today must account for the AI capability level that will exist when the employee is productive—not the level that exists when the requisition is approved.

16.2 The Hiring Treadmill

The acceleration curve creates a compounding problem for GCC workforce planning. Traditional models forecast headcount needs 12–18 months forward based on current workload projections. But AI capability is not static over that planning horizon. A GCC that identifies a need for 50 additional QA engineers based on 2026 workload projections may find that by the time those engineers are hired, onboarded, and productive in mid-2027, AI testing tools have absorbed 40% of the projected workload increase.

The response is not to stop hiring. It is to hire differently. The roles that remain valuable are those where the acceleration curve has not yet reached the capability threshold—complex systems architecture, regulatory judgment, cross-functional stakeholder management, AI governance—and where the capability threshold itself is further away because the work involves dimensions (physical interaction, legal accountability, creative ambiguity) that AI benchmarks do not yet measure.

GCCs that respond to displacement by hiring more people for reskilling may find the reskilling target has moved by the time training is complete. The acceleration curve demands that reskilling programmes target capabilities with longer AI-resistance horizons, not capabilities that are currently just below the automation threshold.

“We approved a headcount plan in January. By July, the AI tools we were evaluating could do 30% of what we were hiring for. By the time the new hires were productive in October, it was closer to 45%. We are planning for a world that no longer exists by the time the plan is executed.”

—GCC Head of Delivery, Technology Sector

17. The Structural Paradox

There is a profound irony at the centre of the GCC model. The entire value proposition—standardised, process-driven, rules-based cognitive work executed at scale by trained professionals—is a near-perfect description of what modern AI systems are engineered to replace.

GCCs succeeded by taking ambiguous, complex work from Western headquarters and converting it into something legible: repeatable processes with documented steps, defined inputs and outputs, quality metrics, and standard operating procedures. They built playbooks. They created SOPs. They standardised outputs and reduced variance. These are the exact preconditions for automation.

Every efficiency gain that made a GCC more attractive to a parent company also made it more automatable. Every process document, every workflow diagram, every quality checklist—these are, in effect, training manuals for AI agents. The industry’s obsession with documentation, process standardisation, and output consistency—the qualities that drove decades of growth—is now the primary accelerant of its own automation.

ORBYS360 STRATEGIC INSIGHT

Every process document ever written for a GCC operation is, in effect, a training manual for an AI agent. The GCC industry's obsession with documentation, process standardisation, and output consistency—virtues that drove decades of growth—is now the primary accelerant of its own automation.

The paradox extends deeper. The management systems that GCCs developed to demonstrate value—detailed activity tracking, output measurement, SLA dashboards, utilisation metrics—have created a comprehensive record of what each role does, how long it takes, and what quality looks like. This record is exactly the data that AI systems need to learn the task, and that parent company executives need to calculate the return on replacing humans with machines.

The paradox also operates at the level of individual professionals. The GCC employee who has spent years perfecting a process—reducing error rates, improving throughput, documenting edge cases—has, in effect, been building the specification for the AI system that will replace her. The more expert she becomes, the more complete her documentation, the easier her job becomes to automate. Her competence is not protection; it is acceleration.

This is not a comfortable insight. It is, however, the structural reality that GCC leadership must understand if they are to navigate the transition honestly. The paradox cannot be resolved by working harder or more efficiently within the existing model. It can only be resolved by changing the model—moving from work that AI can learn to work that requires the judgment, creativity, and contextual intelligence that AI cannot replicate.

18. Enterprise HQ Sentiment: The Bloat Problem

While the GCC industry has been processing the implications of AI, a more immediate threat has accumulated inside the boardrooms of the multinationals that fund GCC operations. Orbys360’s survey of multiple global enterprises reveals a shift in sentiment: from satisfaction to scepticism, and in a meaningful minority of cases, to formal restructuring review.

 ORBYS360 HQ SURVEY FINDING

67% of enterprise CFOs and COOs stated they are actively reviewing the headcount justification of their GCC operations in light of AI capabilities. A further 22% said they had already initiated formal GCC footprint reduction studies. Only 11% reported no active review. Of those conducting reviews, 83% cited AI capability as the primary trigger—but pre-existing concerns about headcount size, accountability gaps, and ROI opacity as the primary content.

Survey Statement	Agree %	Urgency
Our GCC headcount has grown faster than the value it delivers	71%	HIGH
AI can now do at least 40% of what our GCC does, for less cost	78%	HIGH
We are unlikely to grow our GCC headcount over the next 2 years	64%	MED-HIGH
We are actively evaluating whether our GCC should consolidate or reduce	58%	MED-HIGH
Our GCC leadership has not presented a credible AI transition strategy	73%	HIGH
We would not replicate our current GCC structure if building from scratch	69%	HIGH
We lack clear ROI visibility on our GCC operations	54%	MED-HIGH
Cost per unit of output from our GCC has not decreased in 3 years	61%	HIGH
We are piloting AI to replace specific GCC function categories	49%	MED-HIGH
The GCC model is less strategically important than 5 years ago	57%	HIGH

Source: Orbys360 Enterprise HQ Survey, 2025.

18.1 The Three Patterns of GCC Bloat

01 | Managerial Layer Inflation. GCCs carry approximately 1 manager for every 5–6 individual contributors, a ratio that has not compressed as headcount has grown. 61% of HQ respondents cited excessive management layers as a primary cost driver. As GCCs grew, they hired managers to manage the growth, creating a self-reinforcing overhead spiral.

02 | Scope Creep Without Accountability. GCC mandates expanded organically without corresponding accountability for outcomes. HQ executives describe mandate drift—new functions absorbed without old ones retired. 54% of HQ respondents said they lack clear ROI visibility on GCC operations, a remarkable admission about operations costing tens of billions annually.

03 | Hiring as a Signal of Strategic Importance. GCC leadership has historically used headcount growth as a proxy signal of strategic importance. With AI making output-per-head measurable in real time, that proxy is being replaced by direct output measurement—and the results are exposing the inflation. 66% said GCC productivity metrics have not kept pace with headcount growth.

“We built the GCC to reduce cost. Now we are discovering the GCC itself is a cost we need to reduce. The irony is not lost on us.”

—CFO, Global Financial Services Firm

The structural mechanisms that produced these three patterns—and the role archetype that anchors them—are examined in the following section.

19. The Seat Warmer Economy: How GCCs Built Their Own Structural Vulnerability

The 50% structural-exposure figure introduced earlier in this report is not solely a consequence of AI capability advancing beyond expectation. It is, in significant part, a consequence of how GCCs chose to grow. Over two decades, the industry built organisations optimised for activity rather than impact—staffed with roles designed to execute defined tasks, not to own outcomes. The AI reckoning has not created this vulnerability. It has exposed it.

The previous section documented the bloat patterns as HQ executives now perceive them. This section takes the diagnosis one layer deeper—to the role archetype that produces those patterns, the structural mechanisms that built it at scale, the hidden cost structure it carries, and the redesign path that closes it before the reckoning forecloses that option.

The ‘seat warmer’ is the role archetype that defines this vulnerability—and it is not a description of individual failure. The term requires analytical precision before its structural implications can be properly understood.

19.1 Defining the Seat Warmer: Role Archetype, Not Performance Verdict

The term ‘seat warmer’ is frequently misread as a critique of individual capability. This is analytically imprecise and organisationally counterproductive. A seat warmer is a role archetype, not a performance classification. The individuals who occupy such roles may exhibit high process adherence, consistent delivery within defined parameters, and low error rates. Their performance, by the metrics their organisations have chosen to apply, may be entirely satisfactory.

The problem is not with the individual. It is with the organisational design that created a role whose defining characteristics—execution of predefined tasks without contextual ownership, absence of linkage to measurable business outcomes, limited scope for judgment or decision-making, and persistence within the system due to structural rather than strategic logic—make it simultaneously easy to perform, invisible to eliminate, and maximally exposed to AI displacement.

The seat warmer role can be distinguished from value-generating roles along seven analytical dimensions. The following table makes that distinction explicit, drawing on Orbys360's role-level analysis across 200+ GCCs tracked in depth within the global survey universe.

Dimension	Seat Warmer Role	Value-Generating Role
Definition	Executes predefined tasks without contextual ownership or business-outcome linkage. Role persists due to structural design, not strategic need.	Owens a defined outcome, exercises judgment, and can directly articulate the business impact of their work.
Measured by	Activity metrics: tickets closed, transactions processed, hours logged, SLA adherence.	Outcome metrics: quality of decisions, business impact generated, capability advancement.
Scope of judgment	Minimal. Work follows documented rules, escalation matrices, and process scripts.	Significant. Work requires contextual intelligence, domain expertise, and adaptive reasoning.
Linkage to outcomes	Absent or indirect. Output contributes to a chain of tasks; no single role owns the result.	Direct. Role holder is accountable for a measurable result that the organisation values independently.
AI displacement risk	CRITICAL Role characteristics—repetition, rule-following, low contextual dependency—match precisely the tasks AI agents automate first.	LOW-MODERATE Contextual judgment, domain ownership, and outcome accountability are the qualities AI cannot replicate.
Individual performance	Often high by traditional metrics—process adherence, low error rates, consistent delivery. The problem is structural, not personal.	Evaluated on business impact, not process compliance. Performance is intrinsically linked to outcome quality.
Persistence mechanism	Structural inertia, headcount-as-signal logic, lift-and-shift legacy, and fragmented outcome ownership.	Mandate clarity, outcome accountability, and deliberate organisational design.

The most important observation from this comparative framework is the AI displacement risk differential. Seat warmer role characteristics—repetition, rule-based execution, documented process adherence, and low contextual dependency—are precisely the qualities that enable AI automation. The task profiles of seat warmer roles in IT delivery, BPM operations, and engineering support are not incidentally exposed to AI displacement. They are structurally aligned with the capabilities that LLMs, agentic systems, and AI coding tools are designed to absorb.

💡 ORBYS360 KEY FINDING

The seat warmer phenomenon is not the result of poor hiring decisions or individual underperformance. It is the predictable outcome of organisational design choices that prioritised headcount growth, process stability, and activity metrics over outcome accountability, mandate clarity, and value measurement. GCCs did not accidentally create seat warmers. They built systems that systematically produced them.

19.2 Structural Origins: Five Mechanisms That Built the Archetype at Scale

The seat warmer phenomenon did not emerge through a single design failure. It accumulated through five reinforcing structural mechanisms that operated across IT, BPM, and engineering service lines simultaneously—each one rational in its local context, collectively producing an organisational architecture that is now acutely vulnerable to the AI reckoning.

Mechanism	Description and Pillar Examples
01 Legacy of Lift-and-Shift Models	A significant proportion of GCCs originated through lift-and-shift transitions in which existing processes were relocated geographically with minimal redesign. Continuity was prioritised over optimisation; replication over reinvention. Roles were never re-evaluated for relevance after migration. The playbooks that described them were imported wholesale from the originating HQ environment, along with the assumptions—and the structural inefficiencies—those playbooks contained. IT: junior development and QA roles replicated without redesign as AI accelerated delivery cycles. BPM: transaction processing and compliance roles imported with no automation review. Engineering: drafting, documentation, and test-execution support roles copied from captive engineering labs.
02 Persistence of Service Provider Economics	Despite being captive entities, many GCCs continue to operate with the implicit logic of the BPO model they were designed to replace. Headcount is used as a proxy for value. Utilisation rates substitute for outcome measurement. Stability signals success. These metrics incentivise scale without scrutiny, allowing roles to persist regardless of their strategic contribution—because the measurement system was never designed to detect redundancy. IT: ticket volume and code-line metrics reward throughput regardless of business value delivered. BPM: transaction counts and SLA adherence create perverse incentives to process, not to improve. Engineering: drawing counts and test execution hours replace meaningful engineering contribution metrics.
03 Fragmentation of Outcome Ownership	In numerous GCC operating models, strategic intent is retained at headquarters, execution is distributed to the GCC, and accountability for outcomes is diffused across the chain. The result is work that is task-oriented rather than outcome-driven, disconnected from business impact, and resistant to elimination precisely because no one owns the outcome it ostensibly serves. In such environments, activity becomes a substitute for value—and the activity can be evaluated positively even as the outcomes it produces become irrelevant. IT: feature delivery measured without linkage to product usage, customer value, or system performance. BPM: process compliance measured without connection to the business decision the process was designed to support. Engineering: design output measured without connection to product commercialisation or engineering advancement.
04 Leadership Risk Aversion and Structural Inertia	Transforming a GCC from activity-based to outcome-driven entails role rationalisation, capability redesign, and cultural disruption. These are inherently high-friction interventions. GCC leadership that has built its authority on headcount size, delivery scale, and process maturity faces a direct personal risk in sponsoring changes that reduce the first and challenge the second. Consequently, leadership often defaults to preserving stability—thereby reinforcing the structural conditions that created the seat warmer phenomenon in the first place. IT: senior IT leaders resist AI adoption that would reduce the scale of organisations built around FTE delivery. BPM: operations directors protect high-volume process teams whose elimination would reduce their organisational footprint. Engineering: engineering managers preserve specialised sub-teams whose mandates have become peripheral to core product development.

Mechanism	Description and Pillar Examples
05 Headcount as a Signal of Strategic Importance	The third pattern of bloat identified in the previous section—headcount as a signal of strategic importance—is the surface visible expression of a deeper organisational logic. Orbys360’s HQ research documents a pattern that GCC leaders have rarely articulated explicitly: headcount growth has historically been used as a proxy signal of strategic importance to parent company stakeholders. A GCC with 10,000 employees commanded more board attention than one with 4,000. This incentive structure encouraged growth for its own sake—absorbing work that might have been automated, creating roles that could have been eliminated, and expanding mandates without corresponding accountability. With AI making output-per-head measurable in real time, that proxy is being dismantled. The inflation it enabled is being exposed. IT: GCC IT headcount grew to demonstrate capability breadth; AI is now exposing the output-per-head deflation. BPM: BPM team scale signalled operational maturity; throughput-per-agent metrics are now replacing headcount as the measure. Engineering: engineering team size was used to justify engineering GCC investment; innovation and IP output are now the interrogated metrics.



ORBYS360 STRATEGIC INSIGHT

Taken individually, each of these five mechanisms has a plausible institutional logic. Lift-and-shift was a sensible risk-management approach in the early GCC era. Service provider economics provided a familiar measurement framework. Outcome fragmentation reflected genuine complexity in global operating models. Leadership risk aversion is a rational response to high-friction change. Headcount signalling worked—until AI made output-per-head measurable in real time. Taken together, these five mechanisms produced an organisational model that is optimised for the conditions of 2010 and structurally misaligned with the conditions of 2026.

19.3 The Hidden Cost Structure of Seat Warmer Economies

The presence of seat warmer roles across IT, BPM, and engineering GCC functions introduces a set of second-order costs that are systematically underreported in standard GCC performance reviews. These costs do not appear on the balance sheet in the period when the structural conditions are established. They materialise at the moment of reckoning—when AI deployment, HQ scrutiny, or talent attrition force a structural assessment that the organisation is unprepared to make.

Cost Category	Mechanism
01 Strategic Marginalisation	GCCs dominated by seat warmer roles are perceived by parent company executives as cost centres rather than value creators—operationally necessary but strategically replaceable. Orbys360’s HQ research found that 57% of enterprise decision-makers describe the GCC model as less strategically important than five years ago. This perception directly limits GCC influence in enterprise planning, technology investment decisions, and business unit alignment—creating a self-reinforcing cycle in which lack of strategic influence produces more seat warmer roles, which produces less strategic influence.
02 Innovation Suppression	Seat warmer roles, by design, do not incentivise process reimagination, automation, or the elimination of redundant work. In IT, this manifests as maintenance-mode delivery cultures that generate no engineering IP. In BPM, it produces process teams that optimise existing flows rather than question whether the flow should exist. In engineering services, it creates execution-focused teams that accumulate technical debt rather than design capability. Legacy processes persist because the roles that execute them have no incentive—and no mandate—to eliminate themselves.
03 Disproportionate AI Exposure	The defining characteristics of seat warmer roles—repetition, rule-based execution, low contextual dependency, and documented process adherence—align precisely with the task profiles that AI agents automate first, best, and at unlimited scale. A GCC with a high concentration of such roles across IT delivery, BPM operations, and engineering support is not merely under-performing strategically. It is structurally positioned to have the majority of its work displaced by AI in the 2025–2030 window. The seat warmer phenomenon is not a legacy inefficiency that can be corrected before the reckoning arrives. It is one of the primary mechanisms through which the reckoning is expressed.
04 Adverse Talent Dynamics	High-potential professionals—in IT, BPM, and engineering alike—gravitate toward ownership, impact, and growth. In environments dominated by seat warmer roles, such talent exits. The departure is not always visible as a loss because the organisation continues to function at the process level. What leaves is judgment capacity, institutional knowledge, and the latent capability to respond to disruption. What remains is a workforce optimised for compliance rather than capability—precisely the profile least equipped to navigate an AI transition, and most likely to be displaced by the AI systems that are arriving.
05 The HQ Perception Gap	The HQ survey readings reported in the previous section take on a different meaning once the seat warmer phenomenon is introduced as their structural cause. 71% of enterprise CFOs and COOs stating that GCC headcount has grown faster than the value it delivers; 73% saying GCC leadership has not presented a credible AI transition strategy; 61% reporting that cost per unit of output has not decreased in three years—these are not critiques of operational execution. The seat warmers in these organisations are often delivering exactly what they were asked to deliver. They are structural indictments of an organisational design that optimised for activity rather than impact, and that is now visible to every HQ cost analyst for the first time.
06 The Automation Acceleration Trap	The perverse irony of the seat warmer phenomenon is its relationship with AI investment decisions. When HQ decision-makers survey a GCC workforce dominated by rule-based, documented, repetitive roles, they see not a workforce to be developed—they see an automation roadmap. The seat warmer population becomes the primary argument for accelerated AI deployment, reducing the transition runway available to the GCC and the individuals within it. GCCs that have failed to diversify their role composition toward judgment-intensive, outcome-accountable work are, in effect, presenting their parent companies with the case for their own displacement.

The combined effect of these six cost categories is a GCC that has invested heavily in scale—in people, in process, in infrastructure—but that has simultaneously eroded the organisational qualities that make scale defensible: judgment capacity, outcome accountability, innovation culture, and the talent pipeline that sustains all three. When the reckoning arrives—and across IT, BPM, and engineering service lines, Orbys360’s data indicates it has already begun—the GCC finds itself without the internal capability to respond and without the HQ credibility to negotiate the response time it needs.

“The cruelty is not the displacement itself—it is the information asymmetry. Organisations know what is coming. Their employees, in the vast majority of cases, do not.”

—Orbys360 Research Director, 2025

19.4 Structural Alignment with the AI Reckoning

The structural paradox introduced earlier in this report—that GCCs documented their work into training manuals for the AI systems that would replace them—translates, at the role level, into something more specific. The connection between the seat warmer phenomenon and the AI displacement wave is not incidental. It is structural. The defining characteristics of seat warmer roles across IT, BPM, and engineering—task execution without contextual ownership, rule-based performance within defined parameters, documented process adherence, and resistance to elimination through lack of outcome linkage—map with near-perfect precision onto the task profiles that AI agents are engineered to absorb.

Consider the alignment across each service line. In IT, junior development roles that generate boilerplate code, execute documented test scripts, and process Tier 1 support tickets are the seats that AI coding agents fill first—data entry and processing roles in IT and BPM, customer service and query handling roles in BPM, QA and testing engineers across IT and engineering, and junior software developers across IT service lines are the role categories at the leading edge of the displacement curve. In BPM, transaction processing roles that execute predefined compliance checks, reconcile accounts according to documented rules, and route customer queries along scripted decision trees are the seats that agentic AI occupies at 40–60 times human throughput. In engineering, documentation roles, CAD/CAM drafting functions, and test execution support positions that follow defined drawing standards and analysis procedures are the seats that AI design tools are already claiming.

The GCC that has spent two decades building an organisation full of well-designed seat warmer roles has, in effect, produced an extraordinarily detailed specification for the AI systems that will replace them. The process documentation, the standard operating procedures, the SLA frameworks, the training manuals—these are not just operational assets. They are the training data and the task specifications for the AI models that will execute those tasks at a fraction of the cost.

The displacement timeline is not uniform across the three service pillars. BPO and BPM seat warmer roles face the earliest horizon—2025–2027—because voice AI and transaction agent technology are already enterprise-ready. IT seat warmer roles face a 2026–2029 horizon as AI

coding and automated DevOps mature into the enterprise deployment cycle. Engineering seat warmer roles—documentation, drafting, and test execution—face a 2026–2029 horizon as AI design tools and generative CAD platforms reach production-grade capability. Across all three pillars, the direction is identical and the compressed timeline is the same: the displacement of task-execution roles is not a future event. It has begun.

🕒 THE TRANSITION WINDOW

The Structural Paradox at Role Level. The GCC industry spent twenty years turning ambiguous headquarters problems into clean, documented, repeatable processes. It wrote the playbooks. It standardised the outputs across IT delivery, BPM operations, and engineering support. It reduced variance. It was proud of it. It did not realise it was writing training manuals for the machines that would replace it.

19.5 The Path Out: From Seat Warmer Economy to Outcome Organisation

Orbys360's field research across 200+ GCCs tracked in depth identifies a clear pattern: organisations that are successfully navigating the AI reckoning are not those that avoided building seat warmer roles—most built them, because the structural incentives to do so were overwhelming. They are organisations that recognised the seat warmer composition as a structural liability and began redesigning it before HQ scrutiny forced the question.

The window for proactive redesign across IT, BPM, and engineering service lines is approximately 18 months from the point of this publication. After that, the combination of HQ-mandated efficiency reviews, accelerating AI deployment, and talent erosion will make the redesign reactive rather than strategic. The following actions represent Orbys360's recommended sequence for GCC leadership pursuing the transition from a seat warmer economy to an outcome organisation.

Action	What It Means in Practice	IT BPM Engineering
01 Outcome Mapping	Map every role in the GCC against a specific, measurable business outcome—not a task or process. Roles that cannot be linked to a business outcome are candidates for elimination, consolidation, or redesign.	IT: link engineering roles to shipped features, system reliability, or security outcomes. BPM: link process roles to decision quality, cost reduction, or risk mitigation. Engineering: link support roles to product commercialisation, safety performance, or innovation pipeline.
02 Mandate Redesign	Renegotiate the GCC mandate with parent company HQ to explicitly define outcomes rather than activities. Replace SLA-based contracts with value-based frameworks that require the GCC to demonstrate business impact, not process adherence.	IT: shift from ticket-SLA to engineering capability and product contribution metrics. BPM: shift from transaction volumes to decision-support quality and compliance effectiveness. Engineering: shift from drawing counts to IP generation, design-to-manufacture cycle, and failure reduction.
03 Role Reclassification	Systematically reclassify the GCC role portfolio into three categories: roles with direct outcome linkage (retain and develop), roles that can be redesigned toward outcome ownership (invest in transition), and roles that are fundamentally task-execution roles susceptible to AI displacement (managed transition or exit).	Apply the reclassification framework separately across IT, BPM, and engineering functions—the displacement timeline and redesign path differ by service line.
04 Governance Capability Investment	Invest in AI governance, model oversight, and domain-expert judgment capacity. These are the roles that are structurally resistant to AI displacement—not because they are complex to automate, but because they require the contextual ownership and outcome accountability that AI cannot hold.	IT: AI coding agent review, architecture governance, security oversight. BPM: AI model validation, compliance governance, exception handling. Engineering: AI design review, product safety validation, systems integration judgment.
05 Leadership Accountability Reform	Change what GCC leadership is measured on. Leadership that is evaluated on headcount size, utilisation rates, and process stability will always defend seat warmer roles. Leadership evaluated on AI transition progress, value-per-employee metrics, and mandate evolution will eliminate them.	Apply outcome-based leadership metrics uniformly across IT, BPM, and engineering service line heads—not just to the GCC managing director.

The transition from seat warmer economy to outcome organisation is not a headcount reduction exercise. It is a redesign of what the GCC is for—and by extension, of what it is worth to the parent company that funds it. A GCC of 3,000 professionals with clear outcome ownership across IT product governance, BPM intelligence, and engineering IP generation is more valuable—and more defensible against the reckoning—than a GCC of 8,000 executing well-documented tasks that AI will be able to perform before the next annual review cycle completes.

💡 ORBYS360 KEY FINDING

The seat warmer economy was not built through malice or incompetence. It was built through a set of structural incentives—headcount as signal, activity as measure, process as value—that made sense in the world of 2010 and are now liabilities in the world of 2026. The GCCs that escape the reckoning will be those that redesign their role composition with deliberate urgency—across IT, BPM, and engineering simultaneously—before the AI displacement wave and the HQ scrutiny cycle converge on an organisation that is still optimised for a world that no longer exists.

20. The Double Squeeze: AI From Below, Headquarters From Above

What makes the current moment uniquely dangerous for GCC operations is simultaneous pressure from two directions. Previous threats were external—geopolitical instability, currency movements, competitive pressure from rival geographies. The current threat is internal and bilateral. AI is automating the work from below; parent company scrutiny is questioning the model from above.

The current threat cannot be addressed within the existing model because it questions the model itself. AI does not make GCC delivery cheaper; it makes GCC delivery unnecessary for an increasing share of functions. HQ scrutiny does not ask whether the GCC is efficient; it asks whether the GCC is needed.

Among surveyed HQ executives who initiated formal GCC reviews, 83% stated that AI capability was the primary trigger. But the content of those reviews focused on pre-existing concerns: headcount bloat, accountability gaps, strategic alignment, and ROI opacity. AI provided the permission structure for a conversation that many HQ executives had wanted to have for years.

The reinforcing dynamic between the two pressures is worth understanding in detail. When a parent company deploys AI tooling that reduces the need for a particular GCC function, the immediate effect is a reduction in the volume of work flowing to the GCC. The GCC's unit economics worsen because fixed costs (management, infrastructure, facilities) are spread across a smaller volume of work. This worsening shows up in the HQ cost analysis, which strengthens the case for further reduction. The cycle accelerates: AI reduces volume, reduced volume worsens economics, worsened economics justify further AI deployment and headcount reduction.

The cycle is difficult to interrupt from within the GCC. The traditional GCC response to cost pressure—finding efficiencies, reducing attrition, improving utilisation—does not address the fundamental issue, which is that the demand for the GCC's traditional output is declining. Efficiency improvements in a shrinking market are like running faster on a treadmill: the effort is real, but the destination does not change.

Breaking the cycle requires the GCC to change what it delivers, not how it delivers it. That means shifting from output (processed transactions, lines of code, reports generated) to capability (AI governance, domain intelligence, innovation capacity). This shift is the subject of the New GCC Value Proposition section later in this report. What matters for the double squeeze analysis is that the shift must happen while both pressures are active—a degree of difficulty that most GCC leadership teams have not confronted.


ORBYS360 STRATEGIC INSIGHT

GCC leadership must understand that the threat has changed in character. It is no longer external and resolvable through geographic arbitrage. It is internal, structural, and bilateral. The organisations that survive will simultaneously manage the technological disruption from AI and rebuild the value narrative with HQ stakeholders—two tasks requiring very different capabilities and very different conversations.

21. Geographic Impact: India

India ~1.9 million professionals in GCCs **CRITICAL AI RISK**
 GDP exposure: GCC sector $\approx$ 2.1% of national GDP; indirect multiplier estimated at 4.2 $\times$ employment impact. India faces the largest absolute displacement risk globally.

21.0.1 Function Composition and AI Risk



21.0.2 GCC Employment by City

City	Share of GCC Employment
Bengaluru	35%
Hyderabad	20%
Chennai	14%
Pune	12%
Delhi NCR	10%
Other	9%

The concentration of risk is highest in Bengaluru, Hyderabad, and Chennai, where BPO and mid-tier IT delivery dominate. No national AI transition framework exists as of publication.

21.0.3 Key Risk Factors

- No national AI transition framework exists—zero policy infrastructure for displacement mitigation
- Bengaluru alone accounts for 35% of all GCC employment; a single-city systemic risk concentration
- Wage compression and shadow attrition are both active—salary growth has halved while teams shrink via non-backfill rather than announcements
- Urban multiplier effect: Bengaluru’s real estate, retail, and transport are all calibrated to 8–12% annual GCC growth

The urban economics of GCC displacement in India are particularly concentrated. A slowdown to 1–2% growth—let alone a contraction—would have cascading effects on property values, commercial real estate occupancy, and local tax revenues that extend far beyond the IT sector. Hyderabad and Pune face similar dynamics at a smaller scale. Bangalore and Chennai’s GCC base is more weighted toward manufacturing and engineering support functions, which are less exposed, but the BPO and financial services components face the same displacement pressures as their Bengaluru equivalents.



ORBYS360 STRATEGIC INSIGHT

India has the most diversified GCC ecosystem of any country, which provides structural resilience at the top of the value chain. The question is whether diversification into product development and innovation is happening fast enough to absorb the displacement at the BPO and mid-tier IT base.

22. Geographic Impact: Eastern Europe

22.1 Poland


Poland

~320,000 in GCCs

CRITICAL AI RISK

GDP exposure: GCC sector ≈ 1.4% of national GDP; concentrated in Warsaw, Krakow, Wroclaw. Risk is concentrated in financial services operations.

22.1.1 Function Composition and AI Risk



22.1.2 GCC Employment by City

City	Share of GCC Employment
Warsaw	45%
Kraków	28%
Wrocław	20%
Other	7%

Poland’s GCC risk is concentrated in financial services, with 35% of employment in that sector. EU membership creates some regulatory friction around AI deployment through GDPR and the emerging EU AI Act, but does not eliminate the displacement trajectory. The timeline may be stretched by 12–18 months relative to India or the Philippines.

22.1.3 Key Risk Factors

- 35% of Polish GCC employment is in financial services—the highest sector concentration of any geography
- EU AI Act + GDPR create a 12–18 month delay, not a prevention—displacement trajectory unchanged
- CSRD and ESG reporting add scrutiny on headcount decisions; preference for gradual attrition over announced layoffs
- Shadow attrition pattern established: Warsaw, Kraków, Wrocław all showing non-backfill signals
- European parent company reviews triggered by AI capability gap, not headcount—structural review content

 ORBYS360 STRATEGIC INSIGHT

Poland’s EU membership is a buffer, not a shield. The 12–18 month regulatory delay creates a false sense of security. Organisations that use this period to reskill and reposition will survive; those that treat it as a reprieve will face compressed, reactive displacement.

22.2 Romania


Romania

~150,000 in GCC roles

HIGH AI RISK

GDP exposure: GCC sector ≈ 1.1% of national GDP. Key hubs: Bucharest, Cluj-Napoca, Timișoara. Strong in IT delivery and financial services operations.

22.2.1 Function Composition and AI Risk



22.2.2 GCC Employment by City

City	Share of GCC Employment
Bucharest	55%
Cluj-Napoca	28%
Timișoara	10%
Other	7%

Romania has emerged as Central and Eastern Europe’s second-largest GCC hub after Poland, driven by competitive labour costs and a strong technical university pipeline. Software and IT delivery make up 42% of Romanian GCC employment—the highest IT concentration of any major GCC geography—meaning code generation and automated testing are the primary displacement vectors.

22.2.3 Key Risk Factors

- Software and IT delivery at 42% creates heavy exposure to AI coding agents and AIOps tools
- EU AI Act provides a similar 12–18 month delay as Poland; GDPR Article 22 requires human-in-the-loop for automated decisions
- Bucharest–Cluj duopoly creates geographic concentration risk—two-city dependency with limited alternative employment

- Romania's lower wage base vs. Western EU creates some competitive resilience, but the arbitrage logic is eroding with AI
- Growing software product companies in Cluj provide a partial absorbing ecosystem for displaced IT talent



ORBYS360 STRATEGIC INSIGHT

Romania benefits from EU regulatory friction but faces the same structural displacement as non-EU peers. The technical talent base in Cluj in particular represents a genuine re-absorption opportunity—but only if reskilling from delivery roles into product and AI-operations roles is actively funded and structured.

22.3 Czech Republic and Hungary

Czech Republic

~50,000–65,000 in GCCs

High AI risk

Prague and Brno dominant. Strong in finance, automotive, and technology sectors. Higher labour costs than Poland or Romania limit cost-arbitrage advantage.

Hungary

~45,000–60,000 in GCCs

High AI risk

Budapest-centric. Concentrated in financial services and pharmaceutical shared services. Similar risk profile to Poland with smaller absolute exposure.

Both markets face displacement patterns similar to Poland's, with financial services and pharmaceutical shared services dominating the risk profile. Higher labour costs relative to India and the Philippines mean these centres already face competitive pressure from cheaper geographies—a pressure that AI amplifies. The EU regulatory environment (AI Act, GDPR, CSRD) applies equally, creating the same 12–18 month timeline extension.

22.4 Bulgaria

Bulgaria hosts approximately 30,000–40,000 professionals in GCC roles, concentrated in Sofia. The market is cost-competitive and growing in IT and AI-related services, but its smaller scale and later entry into the GCC market mean it is growing into a sector whose foundations are shifting. Bulgaria's lower cost base may provide a brief window of advantage, but the same AI automation that threatens larger hubs will reach Sofia's operations on a similar timeline.

22.5 Regional Pattern

Across Eastern Europe, an estimated 25–35% of GCC roles face medium-to-high automation risk. The EU regulatory framework provides a modest buffer but does not alter the structural trajectory. The region's advantage—proximity to Western European headquarters, cultural alignment, timezone overlap—is being eroded by AI systems that operate without regard to geography, timezone, or cultural context.

23. Geographic Impact: South East Asia

23.1 The Philippines


Philippines

~1.3 million in GCCs and BPO

CRITICAL AI RISK

GDP exposure: GCC/BPO sector $\approx$ 7.5% of national GDP; remittance multiplier creates household-level systemic risk. Highest national concentration of any GCC geography.

23.1.1 Sector Composition and AI Risk



23.1.2 GCC Employment by Metro Area

Metro Area	Share of GCC/BPO Employment
Metro Manila	55%
Cebu	20%
Davao	10%
Clark / Pampanga	9%
Other	6%

Voice-based customer service—the single function most directly in the path of AI voice agent deployment—is the dominant Philippine GCC function, accounting for 40% of sector employment. Combined with back-office BPO at 28%, a full 68% of Philippine GCC employment is in the two highest-risk categories. A contraction of 20–30% in BPO employment would affect the workers themselves and an estimated 4–5 million household dependents who rely on remittance income.

23.1.3 Key Risk Factors

- Voice BPO—the dominant function—is the number one AI voice-agent displacement target globally
- GDP concentration at 7.5% creates macro-level systemic risk extending beyond any single company
- Remittance chain means individual displacement cascades to 4–5 million household dependants

- IBPAP reskilling capacity: 50,000 workers/year vs. hundreds of thousands in the displacement window
- No regulatory framework addresses AI displacement in the GCC/BPO sector as of 2025

The cultural dimension of Philippine GCC displacement is distinct from other geographies. A BPO or GCC job is a social status marker that signifies entry into the professional class. For many workers, it represents the culmination of family investment in education and the foundation of household economic stability. The loss of these positions carries a psychological and social weight that is not captured in economic displacement statistics.

ORBYS360 STRATEGIC INSIGHT

The Philippines faces a policy dilemma with no easy resolution: the BPO/GCC sector was deliberately built over two decades as national economic strategy. Unwinding that dependency while it is still the country's largest private employer requires political will that is difficult to sustain across electoral cycles.

23.2 Vietnam

Vietnam ~40,000–60,000 in GCC/BPO roles **Growing 20–25% YoY**
IT services market \$2.37B (2025), forecast \$3.98B by 2030 (10.98% CAGR). Key cities: Ho Chi Minh City, Hanoi. Approximately 530,000 software developers with 97.5% AI adoption in engineering firms.

Vietnam presents a dual-edged case. The country's IT services sector is growing rapidly, with developer rates of \$20–40/hour (comparable to India) and BPO services running 30–50% cheaper than the Philippines. Significantly lower electricity costs (\$0.084/kWh vs. the Philippines' \$0.15–0.20/kWh) provide operational advantages. The workforce is young, technically literate, and—critically—already heavily AI-augmented, with 97.5% AI adoption across software engineering firms.

The paradox for Vietnam is that it is entering the GCC cycle precisely as AI disrupts the model. The skills and cost advantages that make Vietnam attractive to multinational parent companies are the same attributes that AI is automating globally. Vietnam's window to build scale and diversify up the value chain is shorter than India's was, because the AI displacement wave is already in motion.

23.3 Indonesia

Indonesia ~25,000–40,000 in GCC/BPO roles **Growing 10–15% YoY**
Early-stage GCC market, Jakarta-centric. AI market \$2.4B heading to \$10.88B by 2030. Exported AI services projected from \$230M (2025) to \$750M by 2030. Strong in e-commerce and fintech.

Indonesia’s GCC market is earlier-stage than Vietnam’s but exhibits a striking characteristic: 92% generative AI adoption among knowledge workers, exceeding the global average of 75%. This suggests that Indonesian operations are being built AI-native from the start, rather than retrofitting AI into existing delivery models. The AI services export trajectory—from \$230 million to a projected \$750 million by 2030—positions Indonesia as an AI-augmented delivery hub rather than a traditional labour arbitrage centre.

The risk for Indonesia is that the country is investing in a sector model whose underlying economics are shifting. The opportunity is that, by building AI-native from the outset, Indonesian operations may leapfrog the displacement cycle that is dismantling earlier-generation GCCs.

23.4 Malaysia


Malaysia

~200,000 in GCCs

HIGH AI RISK

GDP exposure: GCC sector ≈ 1.8% of national GDP. Key hubs: Kuala Lumpur, Cyberjaya, Penang. Strong in finance, electronics, and oil & gas shared services.

23.4.1 Function Composition and AI Risk



23.4.2 GCC Employment by Location

Location	Share of GCC Employment
Kuala Lumpur	45%
Cyberjaya	25%
Penang	18%
Johor Bahru	8%
Other	4%

Malaysia is a mature GCC hub with a risk profile similar to Poland’s—concentrated in financial services and manufacturing shared services, with a growing but smaller technology component. Shared services and F&A make up 38% of Malaysian GCC employment—functions fully in the path of agentic AI automation. The country’s higher labour costs relative to the Philippines and Vietnam mean its cost arbitrage advantage is narrower, making AI substitution economics more compelling for parent companies.

23.4.3 Key Risk Factors

- Shared services and F&A at 38% are fully in the agentic AI automation path

- No regulatory framework or national AI displacement policy—unlike Poland, no EU buffer exists
- Displacement pattern mirrors Poland: financial/shared services concentrated in 2–3 cities, non-backfill already visible
- Malaysia’s MSC Malaysia status and government-backed tech zones create growth optics that mask exposure reality
- Growing digital and product function mandates provide a partial upside hedge for skilled workers



ORBYS360 STRATEGIC INSIGHT

Malaysia is in the risk phase where the industry’s growth narrative is still dominant—which is precisely when structural intervention is most effective and least likely to happen. The window for proactive policy and reskilling is 12–18 months before displacement becomes visible in employment statistics.

24. Geographic Impact: Latin America

Latin America is the fastest-growing GCC geography, driven by nearshoring momentum post-2022 and proximity to US parent companies. The region faces a distinct version of the displacement paradox: rapid growth into a sector that is simultaneously being disrupted.

24.1 Mexico

Mexico ~120,000–150,000 in GCCs/SSCs **Growing 15–20% YoY**
 150–180 centres across Monterrey, Guadalajara, and Mexico City. Nearshoring boom driven by US supply chain diversification and timezone alignment.

Mexico’s GCC growth is the fastest in the Americas, fuelled by nearshoring dynamics, USMCA trade alignment, and timezone overlap with US headquarters. Monterrey and Guadalajara have emerged as significant technology hubs, while Mexico City hosts financial services and consulting operations. The nearshoring advantage—same-day collaboration with US teams—is genuine and provides a structural buffer that pure offshore locations lack.

The displacement risk for Mexico is concentrated in the same functional categories as other geographies: BPO, finance operations, and junior IT roles. The nearshoring premium may slow displacement by 12–18 months relative to pure cost-arbitrage centres, because the value proposition includes proximity and timezone alignment alongside cost savings. But AI systems do not require timezone alignment, and the proximity advantage diminishes as more work becomes AI-mediated rather than human-delivered.

24.2 Brazil

Brazil	~80,000–100,000 in GCCs	High AI risk
100–130 centres in São Paulo, Curitiba, and Belo Horizonte. Strong in financial services and manufacturing shared services.		

Brazil's GCC operations serve both US and European parent companies, with a concentration in financial services that mirrors Poland's risk profile. The country's larger domestic market provides some buffer—Brazilian GCCs often serve the local market as well as headquarters—but the back-office and analytical functions that form the core of GCC employment face the same AI displacement pressures as everywhere else. Labour regulations and currency volatility add complexity to restructuring decisions but do not alter the structural trajectory.

24.3 Colombia and Costa Rica

Colombia	~30,000–45,000 in GCCs	Growing 10–15% YoY
40–60 centres in Bogotá and Medellín. Emerging BPO and technology hub with growing bilingual workforce.		

Costa Rica	~40,000–50,000 in GCCs	High AI risk
50–70 centres. Strong in med-tech and financial services (Intel, Amazon, P&G presence). Services outsourcing contributes approximately 1–3% of GDP.		


Both countries are scaling operations in a sector whose traditional economics are being rewritten. Colombia's growth is driven by a bilingual workforce and competitive costs, while Costa Rica has built a reputation for higher-value delivery in med-tech and financial services. The risk for both is that their growth trajectory is based on attracting work that is simultaneously being automated at parent company headquarters.

25. Geographic Impact: Middle East and Africa

25.1 Egypt

Egypt	~120,000 in GCCs	CRITICAL AI RISK
Emerging GCC hub; rapid headcount growth 2019–2024 now at risk of reversal before maturity. Most operations are in BPO and entry-level IT—the highest-risk categories.		

25.1.1 Function Composition and AI Risk



25.1.2 GCC Employment by Location

Location	Share of GCC Employment
Cairo / Smart Village	65%
Alexandria	20%
New Cairo / 6th October	10%
Other	5%

Operations are almost entirely BPO and entry-level IT—the two highest-risk categories, accounting for 75% of Egyptian GCC employment, with no near-term pivot path.

25.1.3 Key Risk Factors

- Egypt attracted GCC investment aggressively from 2019–2024; AI wave arrived before industry diversity was built
- 75% of operations are in BPO and entry-level IT—the two highest-risk categories with no near-term pivot path
- National AI strategy (2021) is adoption-focused; no regulatory framework for displacement management exists
- Growth now at risk of reversal before the sector achieved the scale needed to absorb shock
- No equivalent of India’s diversified ecosystem or Poland’s EU regulatory buffer—Egypt has neither shield



ORBYS360 STRATEGIC INSIGHT

Egypt represents a particularly acute case study in timing risk: a country that successfully executed on a deliberate GCC attraction strategy, only to have the displacement wave arrive precisely at the point of maximum exposure and minimum diversification.

25.2 Morocco

Morocco Ranked **26th** in 2026 Kearney Global Services Location Index **Emerging AI hub**

Leading Maghreb outsourcing destination. National AI plan targets 200,000 AI-trained graduates and 50,000 AI jobs by 2030. IT outsourcing costs 20% below Eastern Europe, 30–40% below France.

Morocco is emerging as a nearshore competitor to both Eastern European and Indian GCC hubs, with a labour-cost index score of 94—nearly matching India and the Philippines at 96. The JAZARI Root AI research centre launched in Rabat in early 2026 anchors a national strategy to position Morocco not merely as a cost-arbitrage destination but as an AI-augmented delivery hub. The implications for established GCC geographies are addressed in the following section on the new geographic competitor.

25.3 Rwanda and East Africa

East Africa is the most nascent GCC-adjacent market but is developing rapidly with explicit AI integration. Rwanda launched Africa’s first AI Scaling Hub in April 2025, backed by \$17.5 million from the Gates Foundation, focused on health, education, and agriculture AI applications. Kenya’s BPO revenues are projected at \$272 million in 2025, with AI-assisted customer experience hubs launching to serve EMEA clients. The broader African BPO market is projected to reach \$14.75 billion by 2033.

The East African model is distinct from established GCC geographies in that it is being built AI-native from inception. Rather than retrofitting AI into existing delivery operations, Rwandan and Kenyan hubs are designed from the outset to operate with AI-assisted workflows. This positions them as potential competitors not just to traditional BPO centres but to any GCC operation that has not yet integrated AI into its delivery model.

26. The New Geographic Competitor: AI-Augmented Labour Markets

The arbitrage model that built GCCs is now being used against them—by the next tier of markets.

India, Poland, and the Philippines built the GCC industry on a cost differential: skilled labour at 25–33% of Western rates. A new generation of competitors is applying the same logic one level down, with AI as the equaliser.

26.1 The Double Arbitrage Threat

The established GCC hubs face pressure from two directions simultaneously. From above: AI automates their work entirely, eliminating the need for human delivery regardless of cost. From below: cheaper markets augmented by AI deliver comparable quality at lower cost. The window in which Bengaluru, Manila, and Warsaw occupy the optimal cost-quality position is narrowing from both directions.

“A competitor GCC in Casablanca operating with AI-assisted workflows at 40% of Bengaluru’s fully-loaded cost is not a distant scenario. It is live and tendering for parent company contracts today.”

—Orbys360, 2026

The mechanism is straightforward. AI coding tools (Copilot, Cursor) close the experience gap between a developer with three years of experience in Ho Chi Minh City and a developer with five years of experience in Bengaluru. The AI handles the skill differential; the cost differential remains. A Vietnamese developer at \$20–40/hour augmented by AI produces output comparable to an Indian developer at \$18–40/hour without AI—but the Vietnamese developer is working in a market where BPO services run 30–50% cheaper than the Philippines and electricity costs are half.

Morocco’s positioning is even more pointed. IT outsourcing costs 20% below Eastern Europe and 30–40% below France, with a national AI plan targeting 200,000 AI-trained graduates by 2030. For European parent companies evaluating their Polish or Romanian GCC operations, Morocco offers a nearer timezone, lower costs, and an AI-augmented workforce being built by national policy.

Indonesia’s 92% generative AI adoption among knowledge workers—above the global average—signals that emerging markets are not waiting to be disrupted. They are building AI-native operations that leapfrog the traditional GCC maturity curve.

Rwanda and East Africa represent the next frontier: the continent’s BPO market projected at \$14.75 billion by 2033, with AI-assisted CX hubs already operational and backed by significant international investment.

ORBYS360 STRATEGIC INSIGHT

The established GCC hubs face a double arbitrage threat. From above: AI automates their work entirely. From below: cheaper markets augmented by AI deliver comparable quality at lower cost. The arbitrage model that built the GCC industry in Bengaluru, Manila, and Warsaw is now being deployed against them by Casablanca, Ho Chi Minh City, Kigali, and Jakarta. The competitive advantage of being the cheapest source of skilled human cognition is no longer a durable position when AI equalises skill and the next tier of markets offers even lower costs.

27. The Labour Economics: Wage Compression, Hiring Freezes, and Shadow Attrition

The visible indicators of GCC displacement—formal layoffs, public restructuring announcements, earnings call disclosures—have so far been limited. This has created a false sense of security. The displacement is happening through mechanisms that are difficult to observe from the outside.

Hiring Freezes and Non-Backfilling. GCC headcount is not being reduced through layoffs; it is being reduced through non-replacement of departed employees. When a team member leaves, the position is not backfilled. The remaining team is expected to absorb the workload, often with AI tools that management has introduced as productivity aids. Hiring volumes at major GCCs dropped 22% year-on-year in 2024, while parent company revenues grew.

Wage Compression. GCC salary growth has flattened across most role categories. In the Indian GCC market, median annual salary increases fell from 9–12% in 2022 to 4–6% in 2025–2026 for junior roles, a decline that exceeds the broader labour market trend and is consistent with reduced bargaining power in a market where supply increasingly exceeds demand.

Shadow Attrition. “Shadow attrition”—Orbys360’s term for the process by which GCC roles are eliminated without formal acknowledgment. A team of twelve becomes a team of nine through natural attrition. The nine are told they are being empowered with AI tools to maintain the same output. No restructuring is announced. No redundancy packages are offered.

The 200+ GCCs in Orbys360’s research show a combined headcount growth rate of just 1.2% in 2025–2026, compared with 8–12% annual growth in the preceding five years. Adjusted for new GCC formations, the underlying headcount change in established GCCs is approximately negative 3%.

The Measurement Problem. None of these mechanisms—hiring freezes, wage compression, shadow attrition—register in the data that policymakers and industry bodies use to track GCC sector health. Nasscom’s annual GCC reports measure aggregate headcount, which continues to grow because new GCC formations (greenfield operations) add headcount to the total even as established GCCs are quietly contracting.

The result is a measurement gap between what the data says and what is happening on the ground. The data says the GCC sector is growing. The reality, within established operations, is that it is contracting—slowly, invisibly, and without the policy attention that visible contraction would trigger.

By the time the measurement gap closes—when aggregate headcount declines become large enough to appear in sector-level statistics—the displacement will be well advanced. Orbys360 estimates that the lag between real workforce contraction and statistical visibility is 18–24 months, driven by the masking effect of new GCC formations and the time required for non-backfilling to translate into net headcount declines at the entity level.

28. The Human Cost

It would be easy to treat the 2.9 million directly exposed professionals as an abstraction. They are not. They are people who made consequential decisions—education investments, housing commitments, family plans—on the basis of a labour market that is being structurally transformed without their knowledge or consent.

Consider a software developer in Bengaluru who spent four years earning a computer science degree and three further years building her skills in a GCC environment. She is 27, holds a mortgage, and her career trajectory looked linear eighteen months ago. Today, her company has announced a productivity initiative and her team of twelve has not backfilled the three who left. She has not been told this is related to AI.

Consider a finance analyst in Warsaw who moved from a regional city to the capital to join a GCC. He is 31, supporting his parents. The job that felt permanent two years ago now has a scheduled efficiency review next quarter. His colleagues have heard that AI tools were trialled in Frankfurt. Nobody has told him what the results were.

Consider a customer service supervisor in Manila who is the first person in her family to hold a professional-class job. She manages a team of twenty. She has been told that a new AI voice system is being tested alongside her team. She has not been told that the test results showed the AI handling 73% of queries without human intervention.

“The cruelty is not the displacement itself. It is the information asymmetry. Organisations know what is coming. Their employees, in the vast majority of cases, do not.”

—Orbys360, 2026

The social effects extend beyond individuals. In Bengaluru, Pune, Hyderabad, and Krakow, the GCC sector has become a major generator of middle-class professional employment. The local economies built around GCC campuses—restaurants, housing, schools, retail—are predicated on a stable and growing workforce. A significant contraction would have macro effects on urban economies that no single company’s ESG report will capture.



THE OBLIGATION

The 50% at risk are not statistics. They are individuals with mortgages, parents, children, and career aspirations built on foundations that are shifting beneath them. The obligation of this industry—of the parent companies that profit from GCC operations, the GCC leaders who manage them, and the governments that depend on them—is to meet that reality with honesty, investment, and urgency.

29. What Survives: The 50% and the Roles AI Cannot Replace

The threat is real, but it is not total. Orbys360's research identifies capabilities that retain structural value. Organisations that survive will restructure around them.

01 | AI Orchestration and Governance. The humans who prompt, train, audit, and govern AI systems are not replaced by them. GCCs that develop AI operations talent will become control centres for enterprise intelligence—managing the machines that manage the work.

02 | Complex Stakeholder Management. Multi-stakeholder negotiations, cross-cultural communication, regulatory navigation, and crisis management require contextual judgment and relationship capital that remain stubbornly human.

03 | Systems Architecture and Design. Senior architects who understand entire enterprise ecosystems—legacy, cloud, and AI layers simultaneously—are more valuable as complexity increases. AI amplifies strategic systems thinking; it does not replace it.

04 | Innovation and Product Discovery. Identifying unmet needs, validating novel ideas, and translating ambiguous business problems into solvable technical challenges remains a human competitive advantage. AI generates variations on existing ideas but does not identify the problems worth solving.

05 | Regulatory and Compliance Intelligence. In banking, healthcare, and defence, the interpretive judgment required for compliance in novel situations requires accountable human professionals. AI outputs alone are not accepted by regulators for high-stakes decisions.

06 | Data Strategy and Ethics Leadership. As AI systems process more enterprise data, the leaders who determine what should be done with that data—questions of ethics, privacy, fairness, and strategic priority—become the most important strategic assets in any GCC.

30. The New GCC Value Proposition: From Cost Arbitrage to Intelligence Hub

The GCC model that delivered value for twenty years was built around a simple arbitrage: the same cognitive work at lower cost. That arbitrage is ending. But the infrastructure, talent base, institutional knowledge, and operational maturity that GCCs have built is not worthless. It is simply no longer sufficient on its own.

Model One: The Intelligence Hub. The intelligence hub repositions the GCC from a delivery organisation to enterprise intelligence infrastructure: a centre where AI capabilities are customised, governed, improved, and amplified by domain-expert human professionals. Instead of delivering work that AI replaces, the hub manages the AI that delivers the work. Instead of being measured on cost per transaction, it is measured on quality of AI output, rate of capability improvement, and strategic value generated.

Several forward-looking organisations are already building toward this model—repurposing finance professionals as AI financial model validators, software engineers as AI coding agent reviewers and trainers, and customer service specialists as conversation designers for AI voice agents.

Model Two: The Innovation Partnership. A smaller segment of GCCs can evolve toward becoming the locus of applied innovation for the parent company in specific domains, markets, or technology areas where geographic proximity, talent concentration, and operational context create real advantages. This model is emerging in India, where some GCCs have taken on product development mandates, built proprietary AI tools, and contributed intellectual property to parent company strategy.

The challenge is that innovation partnership requires a very different talent mix, leadership culture, and organisational mandate than traditional delivery. The transition from a compliance-oriented, process-driven delivery culture to an innovation-oriented, outcome-driven product culture is genuinely difficult and takes years.

The Talent Economics of the Intelligence Hub. The intelligence hub model changes the GCC’s talent economics in ways that are uncomfortable but necessary to confront. A traditional GCC of 5,000 employees might, under the intelligence hub model, employ 1,500–2,000 professionals at significantly higher average skill levels and compensation. The total headcount is smaller. The total compensation bill may be similar or even higher, because the remaining professionals are more skilled and command higher salaries. But the value delivered per professional is dramatically higher, and the work itself is resistant to the next wave of AI capability improvement.

This has implications for the social contract between GCCs and their host cities. A GCC that employed 5,000 mid-level professionals generated significant demand for housing, retail, transport, and food services in its local economy. A GCC that employs 1,800 senior professionals generates less aggregate demand, even if the per-capita spending of those professionals is higher. The transition from the delivery model to the intelligence hub model will, in many cases, produce a net reduction in the economic contribution of GCC operations to their host communities. This is a consequence that honest planning must acknowledge.

🕒 THE TRANSITION WINDOW

Orbys360’s field research suggests that the window for proactive GCC model transition is approximately 18–24 months. After that point, the combination of HQ scrutiny, AI deployment acceleration, and internal talent erosion will make transformative change significantly harder. The organisations that act in 2026–2027 will define what the GCC model looks like in 2030.

The question of who helps GCCs execute this transition—and on what commercial terms—is addressed in the following section on the service provider’s role in the reckoning.

31. The Service Provider in the GCC Reckoning

The GCC reckoning is not a crisis that service providers observe from the outside. It is a crisis arriving at their own door simultaneously—across IT services, Business Process Management (BPM), and engineering services alike. The same AI capabilities hollowing captive GCC workforces are eroding the foundations of the managed services and outsourcing industry at its core. Service providers that approach this moment as an opportunity to pick up displaced GCC volume on the same commercial terms are misreading the landscape entirely.

The reckoning for service providers is structural across all three service pillars. In IT services, AI coding agents are absorbing junior development, QA, and application maintenance at a pace that headcount-billing models cannot survive. In BPM, AI transaction agents process invoices, customer queries, and compliance tasks at 40–60 times human throughput at a fraction of the cost. In engineering services, AI design tools are automating CAD/CAM drafting, test generation, and simulation analysis—the repeatable layer that has historically been the core of offshore engineering delivery.

31.1 The Service Provider's Own Reckoning

The GCC was built on the same premise as the managed services industry: skilled human cognitive labour in lower-cost geographies performing documented, standardised, rules-based work across IT, BPM, and engineering functions. The irony documented throughout this report—that GCCs built themselves to be automated—applies with equal force to the service providers whose business model is identical in its underlying logic. Both industries monetised geographic labour arbitrage. Both are now confronting an AI that performs the same cognitive and technical work without geography.

The displacement dynamic operates through two channels simultaneously across all three service lines. The first is internal: AI tools are reducing the number of human FTEs required to deliver any given scope of IT, BPM, or engineering work—compressing the revenue-per-engagement model and creating deflationary pressure on billing rates across every practice area. The second is external: GCC clients and enterprise buyers are deploying AI internally, reducing their demand for outsourced services in precisely the highest-volume categories—BPM transaction processing, IT application maintenance, QA, junior software development, engineering documentation, and routine analytics.

The combination of internal cost pressure and external demand compression creates a margin vice across IT, BPM, and engineering service lines simultaneously. Service providers that attempt to compete on cost in this environment will find themselves in a race to the bottom against AI systems that have no labour cost floor.

31.1.1 Service Provider AI Disruption Matrix

Service Provider Type	AI Risk	Horizon	Primary Impact Mechanism
Large IT Services (Accenture, Infosys, TCS tier)	HIGH	2026–28	Billable FTE headcount shrinks as AI coding agents absorb junior development, QA, and documentation; clients achieving AI self-sufficiency in software delivery; legacy managed services facing non-renewal
Mid-tier IT Outsourcing & App Maintenance	HIGH	2026–28	Application maintenance and run-the-bank IT work increasingly automated; AI code generation reduces tickets-to-engineer ratio dramatically; pricing power erodes as output-per-FTE increases without billing rate adjustment
BPM / BPO Providers (voice, back-office, F&A)	CRITICAL	2025–27	Voice and transaction processing directly in path of AI agent deployment; AI processes invoices at 40–60× human throughput at sub-cent cost; contract volume collapsing in customer ops, F&A processing, and compliance verification
KPO & Analytics Vendors	CRITICAL	2025–27	Market research, due diligence, competitive intelligence, and report generation reproducible by LLMs; offshore knowledge work value proposition structurally undermined; clients moving to internal AI-augmented teams
HR & Payroll Outsourcing Specialists	HIGH	2026–28	Payroll processing, onboarding administration, and HR operations highly automatable; AI compliance monitoring tools reduce human oversight requirement; workforce analytics increasingly performed by AI platforms
Engineering Services & R&D Support	HIGH	2026–29	AI design tools automate CAD/CAM drafting, simulation setup, and test data analysis; embedded software generation accelerating; but core engineering judgment, systems architecture, and complex failure analysis remain resistant
IT / Management Consulting Boutiques	MODERATE	2027–30	AI augments research, analysis, and benchmarking; but senior advisory, architecture design, change management, and C-suite relationship work remains resistant; firms that embed AI into delivery gain margin; those that do not face commoditisation

31.1.2 The Three Responses—and Their Outcomes

Across IT, BPM, and engineering service lines, Orbys360 observes three distinct strategic postures in response to AI displacement. The posture adopted in the next 12–18 months will determine commercial position through 2030—across all three pillars simultaneously.

Response	What It Looks Like	Commercial Outcome	Trajectory
Denial	Continues to compete on FTE cost across IT delivery, BPM operations, and engineering headcount. Invests in scale and utilisation. Treats AI as a productivity aid within the existing billing model.	Terminal margin compression. Clients achieve AI self-sufficiency. Contract volumes shrink simultaneously across all three service lines. Non-renewal accelerates.	Structural revenue decline. No recovery path within the current model.
Augmentation	Deploys AI tools internally to reduce delivery costs across IT, BPM, and engineering. Maintains labour arbitrage model with AI augmentation. Passes partial savings to clients.	Buys 18–24 months. Revenue per FTE improves but total revenue contracts as headcount falls. Does not resolve the structural problem—only defers it.	Temporary stability followed by accelerated contraction when clients deploy AI independently.
Transformation	Repositions from delivery vendor to transition partner across IT, BPM, and engineering. Builds managed model capability. Assumes outcome accountability. Co-invests in client AI transformation.	Higher-margin, longer-tenure relationships across all three service lines. Defensible against AI commoditisation. Captures the GCC and enterprise transition mandate.	Durable revenue growth. The GCC reckoning becomes a commercial opportunity.

The distinction between Augmentation and Transformation is the most important and most commonly misread across the industry. A service provider deploying AI tools to make existing IT delivery, BPM operations, or engineering headcount cheaper is augmenting—buying time, not building a defensible position. A service provider that redesigns its commercial model around outcome accountability and transition partnership across all three pillars is transforming—and building something that AI cannot replicate, because the value is no longer in the execution of tasks but in the ownership of relationships and accountability for results.

🔍 ORBYS360 KEY FINDING

The service providers most at risk across IT, BPM, and engineering are not those with the lowest margins or smallest scale. They are those that have built the deepest dependence on high-volume, standardised, documented processes—exactly the work that AI automates first, best, and at unlimited scale across all three service lines. The same qualities that made these businesses operationally excellent made them structurally exposed.

31.2 Hybridisation: From Competitor to Architecture Partner

For two decades, service providers and GCCs occupied opposite sides of a historical divide: captive versus outsourced. GCCs were built explicitly to internalise what IT service providers, BPM vendors, and engineering outsourcers had formerly delivered. The AI reckoning is dissolving that opposition. The pressures forcing GCCs to reconfigure create exactly the conditions under which hybrid architectures become not just viable but necessary—and in which the service provider with genuine AI transformation capability becomes an indispensable strategic partner rather than a competing vendor.

Orbys360 identifies three hybridisation archetypes that define the service provider opportunity in the GCC transition. Each applies across IT, BPM, and engineering service lines.

31.2.1 Model One: Managed Capacity Transfer

The service provider takes on IT delivery teams, BPM operations units, or engineering support functions that the GCC is exiting—creating a managed capacity handover rather than an abrupt headcount reduction. The GCC retains strategic oversight and HQ relationship ownership; the SP operates the delivery layer under a defined outcome SLA with contractually bounded transition risk. This model is the highest-volume near-term opportunity across all three service pillars.

- **IT application:** SP absorbs maintenance, QA, and junior development teams being displaced by AI coding agents
- **BPM application:** SP takes on F&A processing, HR operations, and customer ops exiting the GCC perimeter
- **Engineering application:** SP absorbs CAD/CAM drafting, test execution, and documentation workflows being automated
- GCC retains mandate ownership, governance authority, and HQ relationship across all three function types

31.2.2 Model Two: AI-Augmented Hub Co-Investment

The service provider co-invests with the GCC to deploy agentic AI infrastructure across IT, BPM, and engineering functions—reskilling workforces into AI oversight roles and reconfiguring the operating model from delivery to governance. The SP takes outcome-based fees tied to AI performance metrics rather than headcount or activity. Suited to GCCs that have concluded they need to become intelligence hubs but lack the implementation capability to execute the transition across all three service dimensions.

- **IT:** SP co-deploys AI coding platforms, CI/CD automation, and cloud-native architecture; GCC provides domain engineering context
- **BPM:** SP co-deploys AI transaction agents, governance frameworks, and compliance tooling; GCC provides process IP and audit capability
- **Engineering:** SP co-deploys AI design tools, simulation platforms, and test generation; GCC provides domain and product knowledge

- Revenue model: outcome-based fees tied to quality, throughput, and capability improvement metrics—not FTE count

31.2.3 Model Three: The Innovation Spine

The service provider supplies the MLOps infrastructure, rapid prototyping capacity, domain model pre-training, and engineering sprint delivery that enables GCCs—particularly in IT and engineering services—to take on product development and IP generation mandates from parent companies. This model is already emerging in India, where engineering and software GCCs have received product mandates but require AI-native build infrastructure to execute at pace. Revenue shifts from headcount billing to platform licence and IP co-ownership—a structurally superior model for both parties.

- **IT:** SP enables GCC software teams to build AI-native products, contribute proprietary tooling, and generate software IP
- **Engineering:** SP provides AI design acceleration, simulation platforms, and digital twin infrastructure for product engineering mandates
- **BPM:** SP enables GCC process teams to build proprietary AI governance frameworks and compliance intelligence IP for parent company reuse
- Best fit: India-based GCCs in software, industrial, automotive, semiconductor, and life sciences engineering sectors



ORBYS360 STRATEGIC INSIGHT

The service providers that win in the GCC reckoning are not those offering cheaper IT, BPM, or engineering labour to replace displaced workers. They are those offering architecture—the hybrid operating models, AI governance frameworks, and reskilling pipelines that let GCCs transform on their own terms across all three service dimensions. The GCC is not looking for a vendor. It is looking for a transformation partner with capability across IT, BPM, and engineering—and with skin in the game.

31.3 Value Addition Through Managed Models

GCC clients—whether restructuring IT delivery, BPM operations, or engineering support functions—are no longer buying headcount. They are buying outcomes, transition assurance, and capability transfer. The managed model, in which the service provider assumes accountability for results rather than activity, is the commercial architecture through which service providers can build durable, defensible revenue across all three pillars.

31.3.1 Seven Managed Model Plays

#	Play	Scope	Description	Revenue Model	Primary Market
1	AI-Augmented IT Managed Services	IT Services, Cloud, DevOps	SP transitions IT managed service contracts from FTE-billing to outcome-SLA, embedding AI coding agents, automated ops tooling, and intelligent monitoring. Client retains application ownership; SP guarantees uptime, velocity, and cost-per-outcome.	Outcome SLA + platform licence	Enterprise GCCs; IT outsourcing clients
2	Agentic BPM Operations	F&A, HR, Compliance, Customer Ops	SP operates AI agent fleets managing transaction processing, reconciliation, customer query routing, and compliance verification. Human oversight retained for exception handling and regulatory sign-off.	Per-transaction + governance retainer	GCC F&A, HR Ops, Customer Ops
3	Engineering Intelligence Services	R&D Support, Product Eng, Testing, CAD/CAM	SP provides AI-augmented engineering delivery covering automated test generation, AI-assisted CAD/CAM workflows, embedded software acceleration, and simulation-driven design support.	Capability retainer + IP co-ownership	Manufacturing, automotive, aerospace GCCs
4	AI Governance-as-a-Service	All Service Lines, Regulated Sectors	SP deploys and manages AI governance frameworks covering model validation, bias auditing, output quality assurance, and regulatory compliance (EU AI Act, GDPR Article 22, sector standards).	Compliance retainer + audit fees	EU-operated GCCs; financial services, healthcare
5	Workforce Transition Management	IT, BPM, Engineering Reskilling	SP designs and operates reskilling programmes converting professionals into AI operations, orchestration, and oversight roles. Cost: \$3–8K per employee vs. \$40–80M severance for a 5,000-seat GCC.	Per-head reskilling + placement fees	India, Philippines, Egypt
6	Innovation Hub Enablement	Product Eng, IP Generation, MLOps	SP provides MLOps infrastructure, rapid prototyping capacity, domain model pre-training, and sprint delivery enabling GCCs to take on product development and IP generation mandates.	Platform licence + IP co-ownership	India GCCs with engineering depth
7	HQ Mandate Repositioning & Transition Advisory	IT, BPM, Engineering Advisory	SP acts as strategic counsel to GCC leadership in negotiating a new mandate with parent company HQ. Covers AI ROI modelling, board-level narrative development, and functional transition roadmaps.	Advisory retainer + transformation fee	GCCs under HQ review; CFO-driven efficiency programmes

31.3.2 The Revenue Model Shift

Moving from labour arbitrage to managed outcomes requires service providers to restructure their commercial model across IT, BPM, and engineering service lines simultaneously.

From (Legacy Model)		To (Managed Model)
FTE headcount billing (IT / BPM / Engineering)	→	Outcome SLA—velocity, quality, cost-per-unit
Labour arbitrage margin	→	Platform licence + IP co-ownership revenue
Scope-based project delivery	→	Continuous managed capability retainer
Delivery accountability for outputs	→	Transition accountability for business results
Scale through headcount addition	→	Scale through AI agent fleet expansion
Geographic cost advantage as differentiator	→	Governance and domain expertise as differentiator

31.3.3 The Financial Case

The financial argument for managed model positioning is grounded in the displacement cost arithmetic:

- Reskilling cost: \$3,000–8,000 per employee across IT, BPM, and engineering roles for AI operations and governance training
- Severance and restructuring: \$40–80 million for a 5,000-seat GCC undergoing unmanaged displacement
- Managed transition cost: 15–25% of unmanaged displacement when total costs are included—the financial case for service provider engagement is direct and quantifiable
- Cost premium for delay: approximately 30–50% per year of inaction, driven by top talent attrition, deteriorating HQ negotiating position, and compounding institutional knowledge loss

The GCC that invests \$15 million in managed transition across its IT, BPM, and engineering functions—with SP fees included—to avoid \$60 million in severance is a client that sees direct, quantifiable return from the engagement. This is fundamentally different from competing on headcount cost, and it produces a client relationship that is longer-tenure and more defensible against AI commoditisation.

 **ORBYS360 STRATEGIC INSIGHT**

The GCC reckoning creates a \$46 billion addressable transition opportunity for IT, BPM, and engineering service providers—not in replacing displaced workers, but in building the transition architecture that GCC clients cannot construct alone across all three service dimensions. Service providers that reposition from delivery vendors to transformation partners will capture the AI-era GCC relationship. Those that remain in the labour arbitrage business will be disintermediated by the same AI systems already hollowing their clients.

31.4 What GCC Clients Are Now Buying

GCC clients restructuring IT delivery, BPM operations, and engineering support functions share five consistent purchasing criteria that have emerged in the AI transition period.

Transition Assurance. A credible, costed path from current state to intelligence hub model—covering IT, BPM, and engineering functions—with the SP assuming accountability for execution, not merely advisory support. GCC leaders under HQ review must present this plan to parent company boards within six months. The SP that can deliver it across all three service dimensions owns the relationship.

AI Capability Transfer. Access to AI tooling, governance frameworks, and training pipelines that the GCC cannot build internally at the required speed across IT, BPM, and engineering simultaneously. This includes AI coding platforms, BPM transaction agent infrastructure, engineering design automation, domain model pre-training, and compliance architecture for regulated environments.

HQ Credibility Infrastructure. The narrative, data, and roadmap required to reframe the GCC's value proposition with parent company leadership across all three service lines. Of HQ executives surveyed by Orbys360, 73% stated that their GCC had not presented a credible AI transition strategy. The service provider that closes this gap across IT, BPM, and engineering simultaneously operates at the board level.

Reskilling Pipelines. Structured programmes converting IT delivery professionals, BPM operations staff, and engineering support roles into AI operations, orchestration, and oversight capabilities before displacement creates an unmanageable political and reputational crisis. The reskilling challenge is largest in India (IT and BPM), the Philippines (BPM), and emerging engineering hubs in Poland and Romania.

Outcome Accountability. Service providers willing to take risk on results—not merely activity—across IT, BPM, and engineering service lines, and to structure commercial agreements around measurable outcomes. This is the single most powerful differentiator between a managed model partner and a traditional vendor, and it is the hardest capability for a labour-arbitrage provider to credibly acquire.

31.5 Orbys360 Recommendation for Service Providers

The transition window documented throughout this report—approximately 18–24 months before HQ scrutiny, AI deployment, and talent erosion make proactive transformation significantly harder—applies to service providers across IT, BPM, and engineering service lines as directly as it applies to GCC leadership.

CRITICAL ADVISORY

Service providers that invest in the following capabilities in the next 12 months will be positioned to compete for the IT, BPM, and engineering transition mandates that will define the market through 2030.

- AI governance and compliance tooling across IT, BPM, and engineering deployments—particularly for EU AI Act and GDPR-compliant environments in Poland and Romania
- Reskilling programme design and delivery across all three service lines—with validated outcome tracking for AI operations, orchestration, and oversight roles
- Outcome-based commercial structuring capability—the legal, financial, and delivery frameworks for SLA-driven managed models applicable across IT managed services, BPM operations, and engineering support
- HQ engagement methodology—the ability to operate at board level and present transformation roadmaps covering IT, BPM, and engineering service line transitions to parent company executives
- Domain-specific AI implementation expertise in financial services BPM, software IT delivery, and manufacturing/industrial engineering—the three highest-displacement, highest-urgency combinations in the GCC transition portfolio

The service providers that delay this repositioning will find the market bifurcating: a shrinking volume of traditional staff-augmentation relationships across all three pillars, and a growing volume of transformation partnerships owned by providers that moved earlier. The latter category is higher-margin, longer-tenure, and more defensible across every service line. The choice of which category to compete in is being made now—not in 2027.

32. Financial Modelling: The Cost of Inaction Versus the Cost of Transition

32.1 The Cost of Unmanaged Displacement

Severance and restructuring charges for a GCC of 5,000 employees with 40% displacement typically run between \$40–80 million, depending on jurisdiction and contractual obligations.

The hidden costs are larger. Knowledge loss—institutional understanding of processes, systems, and business context—is consistently cited by CTOs as the most damaging long-term consequence. Morale effects on remaining employees create a productivity drag that persists for 12–18 months after a restructuring.

32.2 The Cost of Managed Transition

Reskilling costs for AI operations and governance training run approximately \$3,000–8,000 per employee. For a GCC of 5,000 where 2,000 employees are reskilled, the total investment is \$6–16 million—roughly 10–20% of the cost of unmanaged displacement.

Role redesign requires \$2–5 million for a GCC of this scale. Mandate evolution requires sustained engagement with parent company leadership, which is a time cost rather than a financial cost, but is the single most important determinant of success.

The comparative calculation is not close. Managed transition costs 15–25% of unmanaged displacement when hidden costs are included, and it preserves institutional knowledge, maintains workforce morale, and positions the GCC for the intelligence hub model.

32.3 The Cost of Delay

Every month of delay in initiating a transition programme increases the eventual cost. The dynamic is straightforward: as AI capabilities improve and HQ scrutiny intensifies, the negotiating position of GCC leadership weakens. A GCC that proposes a transition plan in early 2026 does so from a position of proactive leadership. A GCC that proposes the same plan in late 2027 does so from a position of reactive defence against a restructuring mandate that HQ has already drafted.

Orbys360's modelling suggests that the cost premium for delayed transition is approximately 30–50% per year of delay, driven by three factors: higher attrition of top talent (who leave for AI-era employers), reduced HQ willingness to fund transition programmes (as cost-reduction targets have already been committed), and the compounding effect of institutional knowledge loss as experienced professionals depart without being replaced.

The financial case for immediate action is unambiguous. The only organisations for which delay is rational are those that have already concluded that their GCC operations should be wound down rather than transformed—and even in those cases, managed wind-down is cheaper than chaotic contraction.

33. Regulatory and Policy Environment

European Union. The EU AI Act classifies AI systems by risk level and imposes requirements on high-risk applications. For GCC operations in EU member states, the Act creates compliance obligations for AI in employment, financial services, and healthcare. The practical effect is to slow AI deployment in regulated functions but not to prevent it. GDPR Article 22, which gives individuals the right not to be subject to decisions based solely on automated processing, means that GCC functions involving decisions about individuals require human-in-the-loop architectures even when the AI is capable of operating autonomously.

India. India does not have a comprehensive AI regulation framework. The Digital Personal Data Protection Act (2023) addresses data privacy but does not specifically regulate AI deployment or workforce displacement. No national transition framework addressing GCC displacement exists as of publication.

The Philippines. The Philippine government has been more proactive than most. The IBPAP published a workforce transition roadmap in 2024, and DOST has funded reskilling programmes. Current programme capacity reaches approximately 50,000 workers per year, against a potentially displaced population measured in hundreds of thousands.

Egypt. Egypt adopted a national AI strategy in 2021, focused on economic development and technology adoption. There is no specific regulatory framework for AI deployment in outsourcing operations, and no national workforce transition plan.

Malaysia. No specific AI displacement regulation exists. Malaysia's Digital Economy Blueprint addresses technology growth but not workforce transition. No GCC-specific safety net programme is in place.

Romania and Eastern Europe. EU AI Act and GDPR apply as EU member states, providing the same 12–18 month regulatory friction as Poland. No national AI displacement frameworks exist beyond the EU-level instruments. Romania's digital economy strategy (2020) predates the displacement risk era.

Latin America. No GCC-host country in Latin America—Mexico, Brazil, Colombia, or Costa Rica—has implemented AI displacement regulation or workforce transition frameworks for the outsourcing sector. Growth-oriented policy dominates.

Morocco. Morocco's national AI strategy targets 200,000 AI-trained graduates and 50,000 AI jobs by 2030 but is adoption-focused rather than displacement-focused. No workforce transition framework exists.

The Regulatory Gap. Across all surveyed geographies, a common pattern is visible: regulatory frameworks are either absent, underdeveloped, or focused on promoting AI adoption rather than managing its workforce effects. No GCC-heavy country has implemented regulations that specifically address the displacement of outsourcing workers by AI. No country requires companies to conduct AI displacement impact assessments before deploying AI systems that will

reduce GCC headcount. No country mandates transition support for workers displaced by AI in the outsourcing sector.

The regulatory gap is not necessarily an argument for regulation. Premature or poorly designed regulation could slow AI adoption in ways that harm the competitiveness of GCC-host countries without meaningfully protecting workers. But the absence of any regulatory attention to the workforce effects of AI in the GCC sector is a policy failure that will have consequences. The consequences will be most acute in the Philippines, where the GDP concentration in GCC/BPO employment means that unmanaged displacement has macroeconomic implications that extend beyond any individual company or sector.

34. Strategic Recommendations for GCC Leadership

CRITICAL ADVISORY

Organisations that treat this as a 3–5 year planning horizon will find that the market has moved on a 12–18 month reality. The window for proactive transformation is narrow and closing.

Immediate (0–6 months): Commission an honest AI displacement audit. Map every role against current AI capability benchmarks—the specific systems commercially available today. Most GCC leadership teams are operating from outdated threat assessments. The first step is an accurate diagnosis, which requires intellectual honesty that will be uncomfortable.

Near-term (6–12 months): Launch mandatory AI reskilling. The GCC workforce does not need to disappear. It needs to migrate. Professionals with deep domain expertise who acquire AI orchestration and governance skills are extraordinarily valuable. The reskilling required is achievable: domain knowledge is the harder input, and GCC professionals already have it. AI skills can be taught in 3–6 month programmes. The cost is \$3,000–8,000 per employee, compared with \$20,000–40,000 for severance.

Strategic (12–24 months): Present a credible AI transition strategy to HQ. 73% of HQ executives say their GCC has not presented a credible AI transition strategy. GCC leadership that proactively frames the AI transition as an opportunity—presenting a clear vision for the intelligence hub model with concrete milestones—will be positioned far better than those who wait for HQ to impose a restructuring agenda.

Structural (18–36 months): Redefine metrics and value proposition. The metrics by which GCCs are evaluated—cost per transaction, headcount efficiency, SLA adherence—are metrics for a delivery model being superseded. Organisations that succeed will measure AI model quality, innovation output, strategic capability, and governance effectiveness. Changing the metrics requires changing the conversation with HQ.

34.1 The Leadership Competency Gap

Orbys360's interviews revealed a recurring concern: the leadership competencies required to manage a GCC through the AI transition are materially different from the competencies that produced success in the delivery era. The GCC managing director who rose through operational excellence—who built her career on process optimisation, SLA management, and headcount growth—is now being asked to lead a strategic transformation that requires technology vision, change management, and the ability to sell a new value proposition to sceptical HQ stakeholders.

This is not a criticism of current GCC leadership. It is a structural observation about the mismatch between the skills that the old model rewarded and the skills that the new model demands. Organisations that recognise this gap early and invest in leadership development—or bring in supplementary expertise—will transition more effectively than those that expect the same leaders who optimised the delivery model to reinvent it.

The most effective transitions observed in our case study data involved leadership teams that combined deep operational knowledge (understanding what the GCC does and why) with external AI strategy expertise (understanding what AI can do and how to implement it). Neither capability alone is sufficient.

35. Strategic Recommendations for Enterprise Headquarters

Transparency as a Strategic Requirement. The information asymmetry between what parent companies know about AI displacement and what GCC employees know is an ethical problem and a strategic risk. The most capable professionals—those most valuable for the intelligence hub model—are also most likely to leave when they sense that their employer is withholding information. Transparent communication is the foundation of a managed transition.

Transition Investment Alongside Cost Reduction. The \$46 billion in annual labour cost under AI threat represents a cost reduction opportunity. It also represents a transition investment requirement. Orbys360 recommends allocating a minimum of 15–20% of AI-driven cost savings to workforce transition programmes. Organisations that deploy savings exclusively into AI infrastructure will realise short-term cost benefits and long-term capability deficits.

Governance Frameworks for AI-Human Workforce Architecture. Who is accountable when an AI system produces an incorrect financial report? How are errors detected and corrected? What human oversight is required? How is compliance maintained when work is performed by machines? These are questions most organisations have not answered. Building the governance architecture is a board-level responsibility.

Avoiding the Capability Cliff. The most dangerous pattern in AI-driven GCC restructuring is what Orbys360 terms the “capability cliff”: a rapid headcount reduction that delivers short-term cost savings while destroying the institutional knowledge base required to operate and govern AI systems effectively. The capability cliff typically becomes visible 12–18 months after a major restructuring, when the organisation discovers that it cannot maintain AI output quality,

cannot troubleshoot errors, and cannot adapt its AI systems to changing business requirements—because the domain experts who understood those requirements have been displaced.

The capability cliff is avoidable, but only if the restructuring is designed with governance capacity in mind. Orbys360 recommends that any GCC headcount reduction plan include a mandatory AI governance capacity assessment: a calculation of the minimum number of domain-expert humans required to maintain, audit, and improve the AI systems that are replacing the displaced workforce. That number is typically 15–25% of the displaced headcount, and it represents the irreducible minimum below which cost savings become false economies.

Boards that demand cost reduction targets without requiring governance capacity assessments are accepting a risk that they have not properly quantified. The risk is not that the AI systems will fail immediately. It is that they will degrade gradually, producing outputs that are slightly wrong in ways that are difficult to detect without the domain expertise that has been removed. By the time the degradation becomes visible, the cost of rebuilding the expertise base is substantially higher than the cost of retaining it.

36. Strategic Recommendations for National Governments

National AI Workforce Transition Strategies. Every GCC-host country—India, Poland, the Philippines, Egypt, Romania, Malaysia, Vietnam, Mexico, and the emerging hubs—needs AI workforce transition strategies with the same urgency they applied to attracting GCC investment. Those strategies must include retraining infrastructure at a scale of tens of thousands of places per year, safety nets for displaced mid-career professionals, and policy incentives for GCCs that invest in local workforce transformation.

Honest Economic Forecasting. GCC employment forecasts in most national economic plans have not been updated to reflect AI displacement evidence. This has direct implications for infrastructure investment, education planning, and social safety net design. Governments that plan on the basis of outdated GCC growth projections will be managing a crisis without the tools to address it.

Sector Diversification. For countries with the highest GDP concentration in GCC employment—the Philippines in particular—sector diversification is an immediate economic security requirement. If voice-based BPO employment contracts by 30–40% over five years, the Philippines needs alternative employment sectors capable of absorbing hundreds of thousands of displaced professionals. Building those sectors takes time, and the time to start was yesterday.

The Social Contract Question. There is a deeper question that this report has deliberately avoided answering, because it is a question for societies, not for analysts. The GCC model was, in effect, a social contract: multinational corporations received cheaper labour; host countries received employment, tax revenue, and the transfer of professional skills; and individual workers received salaries, career paths, and social mobility. AI is disrupting one party's obligations under that contract while leaving the others' expectations intact.

Workers who invested years in education and career development on the reasonable expectation that the GCC sector would continue to employ them are now discovering that the other parties to the contract—the multinationals and, in some cases, the governments—have access to information about the future of that contract that the workers do not. The information asymmetry documented in this report is not merely an operational problem. It is a breach of the implicit social contract on which the GCC model was built.

Whether governments choose to address this breach through regulation, through transition investment, through social safety nets, or through some combination is a political question. What is clear is that ignoring the question—treating the displacement as a private matter between employers and employees, to be resolved through market mechanisms alone—will produce concentrated economic hardship in communities that lack the resources to absorb it. The consequences of that hardship will not remain confined to the affected communities. They will become, in time, political and social instability that affects everyone.

37. About This Research

This report is the product of an ongoing research programme. Orbys Research engages with multiple global enterprises operating GCC or shared services operations, draws on senior executive perspectives across the GCC ecosystem, conducts in-depth workforce tracking of 200+ GCCs across major host geographies, and runs a structured expert interview programme covering managing directors, enterprise technology leaders, workforce economists, and AI deployment practitioners. The displacement model draws on published automation research.

The term “displacement” in this report refers to structural reduction in the number of human professionals required to deliver a given volume of output—not necessarily to immediate job losses. Full methodology, data sources, and modelling assumptions are documented in Section 38.

38. Research Methodology, Modelling Assumptions, and Limitations

38.1 Primary Research Sources

This report is the output of an ongoing research programme. Orbys Research engages with multiple global enterprises operating GCC or shared services operations and applies the following primary research instruments across that universe.

Orbys Research GCC Workforce Intelligence Database: Ongoing longitudinal tracking of 200+ GCCs across major host geographies in Asia-Pacific, Eastern Europe, Latin America, and EMEA.

Orbys Research Enterprise HQ Survey: Surveyed senior executives at multiple global enterprises operating GCCs. Findings reflect the views of a substantive cross-section of decision-makers responsible for GCC strategy.

Orbys Research Expert Interview Series: Structured interviews with GCC managing directors, enterprise technology leaders, workforce economists, and AI deployment practitioners.

38.2 Secondary Sources

Oxford Future of Employment Study (Frey & Osborne); McKinsey Global Institute; WEF Future of Jobs Report; Goldman Sachs; Nasscom GCC Report; Everest Group GCC Market Update; ILO “AI and Work”.

38.3 Limitations

AI capability assumptions reflect commercially deployed technology as observed at the time of writing. They do not account for regulatory interventions that may slow enterprise AI deployment in specific jurisdictions, potential capability plateaus, or possible accelerations beyond current trajectories.

Workforce-volume figures are estimates derived from public sector and industry-association data combined with our own analysis. Geographic GDP-impact figures rely on published national statistics and standard sector employment multipliers; readers should treat these as directional indicators rather than precise forecasts.

The pace at which structural displacement translates into actual headcount reduction depends on attrition rates, contract structures, regulatory requirements, and corporate decision-making timelines. All quantitative outputs in this report are estimates and should be read as directional indicators of structural exposure rather than point forecasts.

39. Appendices

Appendix A: Notes on Methodology. Where this report quantifies displacement risk, the underlying scoring is a composite drawing on published automation research from Oxford, McKinsey, the World Economic Forum, Goldman Sachs, the ILO, and others, combined with Orbys Research’s own role-level assessment of currently deployed commercial AI capabilities. Workforce coverage spans the major functional domains found in modern GCCs — business process operations, finance and accounting, software development, quality assurance, infrastructure and operations, analytics and research, human resources operations, and procurement and supply chain. Geographic exposure figures use sector employment multipliers consistent with widely cited industry-association methodologies. All quantitative outputs are estimates and should be read as directional indicators of structural exposure rather than precise forecasts.

Appendix B: Function Threat Matrix — Activity-Level Breakdown.

The following tables present the complete activity-level AI displacement assessment from the Orbys Research GCC Function Threat Matrix. Each activity is assessed for risk level, displacement horizon, primary AI mechanism, and residual human value. Source: Orbys Research, drawing on published automation research from Oxford, McKinsey, WEF, Goldman Sachs, and an expert panel.

B.1 — BPM: Business Process Management

Sub-Function	Activity / Role	Risk	Horizon	AI Mechanism	Residual Value	Human
Finance & Accounting	AP Processing	CRITICAL	2025–26	Invoice AI, OCR+NLP end-to-end	Vendor escalation, exception judgment	
	AR & Collections	CRITICAL	2025–27	AI cash application, predictive collections, auto-dunning	Complex dispute resolution	
	GL & Month-End Close	CRITICAL	2026–27	Automated journal entries, AI reconciliation	Judgmental accruals, audit sign-off	
	Financial Reporting & MIS	HIGH	2026–28	Auto-report generation, narrative AI, BI agents	Strategic commentary, board narrative	
	Tax Compliance & Filing	HIGH	2026–28	RegTech AI, auto-filing, tax rule engines	Novel tax positions, jurisdiction nuance	
	Treasury & Cash Mgmt	HIGH	2027–29	AI cash forecasting, automated bank reconciliation	Counterparty relationships, policy decisions	
	Audit Support & Documentation	HIGH	2026–27	AI evidence gathering, automated workpapers	Auditor judgment, risk assessment	
	Intercompany & Transfer Pricing	HIGH	2026–28	AI intercompany matching, automated TP docs	Cross-border tax strategy	
HR Operations	Payroll Processing	CRITICAL	2025–27	End-to-end AI payroll, statutory computation bots	Legal exceptions, employee escalations	
	Recruitment Ops & Screening	CRITICAL	2026–28	AI resume screening, chatbot scheduling, assessment scoring	Cultural fit, senior-level hiring judgment	
	HR Data Entry & Employee Master	CRITICAL	2025–26	AI-driven HRIS updates, OCR onboarding docs	Sensitivity handling, data governance	
	Benefits Administration	CRITICAL	2026–27	AI eligibility verification, automated enrollment	Complex life events, plan negotiation	
	L&D Administration	HIGH	2026–28	AI content curation, LMS automation, AI training materials	Learning design, cultural adaptation	
	Performance Mgmt Support	HIGH	2027–29	AI feedback synthesis, automated goal-tracking	Manager coaching, sensitive conversations	
	HR Helpdesk & Policy Queries	CRITICAL	2025–27	HR chatbots, policy LLM agents (Tier 1–2)	Complex grievances, disciplinary matters	
	Exit Mgmt & Offboarding	HIGH	2026–28	Automated offboarding workflows, AI checklists	Sensitive HR conversations, legal risk	
Customer Operations	Voice Support (Tier 1)	CRITICAL	2025–26	AI voice agents (GPT-4o/Gemini Live) end-to-end	Emotionally complex, VIP escalations	
	Chat & Email (Tier 1)	CRITICAL	2025–26	LLM-powered chat with CRM, auto-response AI	Regulatory complaints, critical clients	
	Complaint & Dispute Mgmt	HIGH	2026–28	AI complaint classification, resolution engines	Regulatory adjudication, liability decisions	

Sub-Function	Activity / Role		Risk	Horizon	AI Mechanism	Residual Value	Human
	Customer Onboarding & KYC		CRITICAL	2026–27	Document AI for KYC, automated identity verification	High-risk customer judgment, STR filing	
	Customer Data & CRM Ops		CRITICAL	2025–27	AI data enrichment, deduplication, CRM auto-update	Data governance, consent management	
	Tier 2 Technical Support		HIGH	2026–28	AI diagnostic agents, knowledge-base retrieval	Complex multi-system issues	
	Collections & Payments		CRITICAL	2026–27	AI collection agents, predictive dialler	Legal escalation, hardship assessment	
Procurement & Supply Chain	PO Processing		CRITICAL	2025–26	Autonomous PO creation, supplier matching AI	Strategic supplier decisions	
	Invoice & 3-Way Match		CRITICAL	2025–26	End-to-end invoice AI with PO/GR matching	Disputed invoices, fraud exceptions	
	Vendor Master Data		CRITICAL	2025–27	AI vendor validation, deduplication, auto-onboarding	Vendor relationship governance	
	Contract Admin & Tracking		HIGH	2026–28	Contract AI, milestone tracking, clause extraction	Renegotiation, legal interpretation	
	Spend Analytics		CRITICAL	2026–27	AI spend categorisation, anomaly detection	Strategic sourcing recommendations	
	Logistics & Freight		HIGH	2026–29	AI freight booking, route optimisation	Crisis logistics, relationship management	
	Inventory & Demand Planning		HIGH	2027–29	AI demand forecasting, automated replenishment	New product planning, strategic buffers	
Compliance & Risk	KYC/AML Processing		CRITICAL	2025–27	Document AI, identity verification, watchlist screening	Complex risk decisions, STRs	
	Regulatory Reporting		HIGH	2026–28	RegTech AI auto-report generation & submission	Novel regulatory interpretations	
	Sanctions & PEP Screening		CRITICAL	2025–26	Automated real-time screening, AI false-positive reduction	Complex cases, regulatory dialogue	
	Audit Trail & Documentation		CRITICAL	2026–27	AI-generated audit logs, automated evidence	Sign-off, external auditor liaison	
	Data Privacy & GDPR Ops		HIGH	2026–28	AI data discovery, automated DSAR, consent tracking	Novel privacy interpretations	
	Internal Controls Testing		HIGH	2026–28	AI control testing scripts, automated sampling	Control design, risk appetite decisions	
Legal Operations	Contract Drafting (Standard)		CRITICAL	2025–27	LLM contract generation from playbooks	Novel commercial terms, negotiation	
	Contract Abstraction		CRITICAL	2025–26	Document AI clause extraction, obligation mapping	Ambiguous/multi-jurisdiction contracts	
	Legal Research Support		HIGH	2026–27	LLM legal research, auto-summarisation	Strategic legal advice, privilege	
	NDA & Standard Agreement Review		CRITICAL	2025–26	AI contract review, redline generation	Unusual risk positions, board contracts	
	IP & Prior Art Search		CRITICAL	2026–27	AI patent search, prior art mining, FTO analysis	IP strategy, inventor interaction	
Data Entry & Doc Mgmt	Data Entry, Indexing & Coding		CRITICAL	2025	End-to-end document AI; zero human-in-the-loop	Disputed/handwritten edge cases	
	Document Scanning & OCR		CRITICAL	2025	Multimodal AI OCR, near-perfect accuracy	Degraded/rare-language documents	
	Records Mgmt & Archiving		CRITICAL	2026–27	AI auto-classification, retention enforcement	Legal hold decisions, archival strategy	

Sub-Function	Activity / Role	Risk	Horizon	AI Mechanism	Residual Value	Human
	Forms Processing	CRITICAL	2025–26	Forms AI (IDP) across all form types	Unstructured or dis-	dis-
	Technical Doc Writing (Ops)	CRITICAL	2025–26	LLM auto-generation of SOPs, manuals	Sensitive/regulatory tone review	

B.2 — IT Activities

IT Cluster	Activity / Role	Risk	Horizon	AI Mechanism	Residual Value	Human
App Dev: Re-requirements	BA – Requirements Gathering	HIGH	2027–29	LLM user story generation, AI requirements mining	Stakeholder strategy, value prioritisation	
	Functional Spec Writing	HIGH	2026–27	GPT-4 class LLMs drafting specs from transcripts	Domain ambiguity resolution	
	UAT Scripting	HIGH	2026–28	AI test case generation from requirements	Edge case judgment, stakeholder sign-off	
App UI/UX	Wireframing & Prototyping	MEDIUM	2027–29	AI design tools (Figma AI, Galileo)	Creative direction, brand judgment	
	UX Research (Ops)	MEDIUM	2027–30	AI synthetic user personas, automated usability	Nuanced human insight, empathy research	
	Design System Docs	CRITICAL	2026–27	LLM auto-generation of design spec docs	Strategic design decisions	
App Front-End	Junior FE Dev (HTML/CSS/JS)	HIGH	2026–28	AI code gen (Copilot, Cursor) for UI components	Complex interaction design	
	Mid-Level FE Dev	MEDIUM	2027–30	AI-accelerated dev; human for complex logic	Architecture decisions, performance	
	Accessibility & Cross-Browser	HIGH	2026–28	Automated accessibility AI, browser simulation	Edge case & assistive tech nuance	
App Back-End	Junior BE Dev	HIGH	2026–28	AI code gen for CRUD, API scaffolding	Complex business logic, security design	
	Mid-Level BE Dev	MEDIUM	2028–30	AI co-pilot; human drives architecture	Distributed systems, complex data models	
	API Integration	HIGH	2027–29	AI-generated API adapters, integration mapping	Complex partner contracts	
	Microservices Dev (Jr)	MEDIUM	2027–30	AI scaffolding, K8s config generation	Architecture governance, service strategy	
App Dev: QA	Manual QA Tester	CRITICAL	2026–27	AI visual regression, exploratory testing agents	High-risk edge case validation	
	Test Automation Eng (Jr-Mid)	HIGH	2026–28	AI auto-generating test scripts from user flows	Test strategy, coverage architecture	
	Perf & Load Testing	HIGH	2027–29	AI synthetic load, automated bottleneck ID	Capacity strategy, SLA negotiation	
	Security Testing (VAPT, Jr)	MEDIUM	2027–29	AI vuln scanning, automated pen testing	Novel attack surface, risk judgment	
App Dev: DevOps	CI/CD Pipeline Eng (Jr-Mid)	MEDIUM	2027–29	AI-generated pipeline configs	Multi-cloud strategy, governance	
	Release Coordinator	HIGH	2026–28	AI release orchestration, impact analysis	Business sign-off, risk decisions	
	Container & K8s Ops (Jr)	MEDIUM	2027–30	AI-generated Helm charts, auto-scaling config	Platform architecture decisions	

IT Cluster	Activity / Role	Risk	Horizon	AI Mechanism	Residual Value	Human
App Dev: Docs	Tech Doc Writer	CRITICAL	2025–26	LLM auto-gen from code/API specs (Mintlify, Swimm)	Narrative clarity, external-facing tone	
	API Doc Specialist	CRITICAL	2025–26	AI auto-generating OpenAPI/Swagger docs	Developer experience strategy	
Infra: Network	NOC Analyst L1	HIGH	2026–27	AIOps alerting, self-healing network scripts	Novel outage patterns, vendor escalation	
	Network Config Eng (Routine)	HIGH	2027–29	AI config generation, intent-based networking	Complex multi-vendor architectures	
	Bandwidth & Capacity Monitor	HIGH	2026–28	AI capacity analytics, threshold management	Strategic capacity planning	
Infra: Cloud	Cloud Infra Eng (Jr-Mid)	MEDIUM	2027–30	AI IaC gen (Terraform), cloud config assistants	Architecture design, cost strategy	
	FinOps / Cloud Cost Analyst	HIGH	2026–28	AI spend anomaly detection, auto-rightsizing	Business trade-off decisions	
	Cloud Monitoring (Routine)	HIGH	2026–27	AIOps (Dynatrace, Datadog AI) autonomous remediation	Complex root cause, business triage	
Infra: Server & DC	SysAdmin (Routine Tasks)	MEDIUM	2027–29	Automated patching, AI config drift detection	Change governance, vendor negotiation	
	DC Operations Technician	MEDIUM	2028–31	AI predictive maintenance, robot-assisted ops	Physical intervention, hardware decisions	
	Backup & Recovery Ops	HIGH	2026–28	Automated backup validation, AI recovery testing	Business continuity strategy	
Infra: DBA	DBA – Routine Ops	HIGH	2027–29	AI query optimisation, automated perf tuning	Schema design, data model strategy	
	DB Reporting & Extract Ops	HIGH	2026–27	AI SQL generation, auto-scheduled extracts	Cross-domain data governance	
Infra: SOC	SOC Analyst L1 – Alert Triage	CRITICAL	2026–27	AI-SIEM autonomous triage, SOAR automation	Novel threat actor tactics	
	SOC Analyst L2 – Investigation	MEDIUM	2027–29	AI threat hunting, automated indicator correlation	Advanced persistent threat response	
	Vuln Management Analyst	HIGH	2026–28	AI vuln prioritisation, automated patch sequencing	Business risk trade-off decisions	
Infra: ITAM	IT Asset Management	CRITICAL	2026–27	AI discovery, auto-inventory, lifecycle automation	Procurement strategy, licensing	
	Config & CMDB Management	HIGH	2026–28	AI auto-discovery, CMDB population agents	Service dependency governance	
IT Support: Tier 1	Helpdesk – Password & Access	CRITICAL	2025	Full ITSM bot automation; zero human touch	Accessibility/complex identity edges	
	Helpdesk – Software & FAQs	CRITICAL	2025–26	AI chat agents, system integration	Unusual/unapproved software requests	
	ITSM Ticket Logging	CRITICAL	2025–26	AI auto-classification, priority, routing	Ambiguous/high-priority escalations	
IT Support: Tier 2	App Support Analyst	HIGH	2026–28	AI knowledge retrieval, auto-runbook execution	Complex cross-app integration	
	Middleware Support	MEDIUM	2027–29	AI root cause analysis, automated retry	Business-critical integration design	
	Reporting & Dashboarding	CRITICAL	2026–27	AI-generated reports, self-service BI via LLMs	Executive narrative, data strategy	
IT Support: ITSM	Incident Mgmt Coordinator	HIGH	2027–28	AI incident orchestration, impact prediction	Major incident communication	

IT Cluster	Activity / Role	Risk	Horizon	AI Mechanism	Residual Value	Human
	Problem Management Analyst	MEDIUM	2027–29	AI root cause clustering, pattern analysis	Strategic remediation decisions	
	Change & Release Advisory	MEDIUM	2027–30	AI change risk scoring, conflict detection	Business impact assessment	
IT Support: End-User	End-User Training	CRITICAL	2026–27	AI training content, LLM learning assistants	Change management, cultural adoption	
	IT Procurement Ops (Routine)	HIGH	2026–28	AI purchase request processing, catalog automation	Strategic vendor relationship	
Data & Analytics	ETL / Pipeline Dev (Jr-Mid)	HIGH	2027–29	AI code gen for pipelines, auto-schema mapping	Complex data architecture decisions	
	Data Quality Analyst	HIGH	2026–28	AI data quality profiling, anomaly detection	Business data governance policy	
	Data Catalogue & Metadata	HIGH	2026–27	AI auto-metadata tagging, lineage generation	Enterprise data strategy	
	BI Report Developer	HIGH	2026–27	AI-generated dashboards, NL BI queries	Strategic analytics narrative	
	Data Analyst (Routine)	HIGH	2026–28	LLM data analysis agents, auto-insight gen	Hypothesis framing, business context	
	Data Warehousing (Jr)	MEDIUM	2027–29	AI schema design suggestions, DW migration tools	Dimensional strategy, governance	

B.3 — ER&D: Engineering Research & Development

ER&D Cluster	Activity / Role	Risk	Horizon	AI Mechanism	Residual Value	Human
Product Design	Concept Research & Benchmarking	HIGH	2026–28	AI literature mining, competitive synthesis	Creative direction, customer empathy	
	DFM Review	MEDIUM	2027–29	Generative AI DFM analysis, constraint scanning	Supplier partnership, cost trade-offs	
	Eng Change Mgmt (ECO/ECR)	HIGH	2026–28	AI impact analysis, automated ECO routing	Cross-functional change negotiation	
	2D Drafting & Detailing	CRITICAL	2026–27	Generative CAD AI auto-generates 2D from 3D	Customer specification interpretation	
	3D CAD Modelling (Jr-Mid)	HIGH	2027–29	AI-assisted parametric modelling, geometry gen	Complex assembly design, innovation	
	CAD Data & PLM Admin	HIGH	2026–28	AI PLM cleansing, BOM validation, release workflows	Configuration governance	
	BOM Mgmt & Maintenance	HIGH	2026–27	AI BOM extraction, automated part substitution	Sourcing strategy, cost trade-offs	
	Technical Data Package (TDP)	HIGH	2026–27	AI TDP assembly, automated spec compilation	Customer review, contractual accuracy	
	Standards & Regulatory Checking	HIGH	2027–28	AI standards mapping against design parameters	Novel regulatory scenarios	
Simulation	Linear Static & Stress (FEA)	HIGH	2027–28	AI-driven FEA setup, mesh automation, auto-interpret	Failure mode judgment, design iteration	
	Fatigue & Durability	MEDIUM	2027–29	AI load case generation, SN curve correlation	Material uncertainty, field failure data	
	Topology Optimisation	MEDIUM	2028–30	AI-augmented topology tools, generative geometry	Manufacturing feasibility, cost balance	

ER&D Cluster	Activity / Role	Risk	Horizon	AI Mechanism	Residual Value	Human
	Routine CFD (External Aero)	HIGH	2027–29	Physics-informed ML replacing full CFD	Novel flow regimes, accreditation	
	Thermal & Heat Transfer	MEDIUM	2027–29	AI surrogate models for thermal prediction	Complex multi-physics, safety-critical	
	CFD Post-Processing	HIGH	2026–27	AI auto-report from simulation results	Engineering judgment on outliers	
	EMC Pre-Compliance	HIGH	2027–28	AI EMC simulation, automated regulatory checklists	Test lab coordination, authority submissions	
	Control System Modelling (Jr)	MEDIUM	2028–30	AI-assisted model gen, automated parameter tuning	System stability, safety certification	
	Multi-Body Dynamics	MEDIUM	2028–30	AI surrogate models for standard MBD	Novel mechanism design, certification	
Testing & Validation	Test Plan Prep & Scheduling	HIGH	2026–28	AI test planning from requirements, auto-scheduling	Risk prioritisation, lab negotiation	
	Test Data Acquisition	HIGH	2026–27	AI sensor data processing, anomaly detection	Experimental judgment, calibration	
	Test Report Writing (Standard)	CRITICAL	2025–27	LLM auto-generation from raw data & templates	Certification sign-off, customer narrative	
	Reliability & MTBF	HIGH	2027–28	AI Weibull analysis, field data correlation	Warranty strategy, field risk decisions	
	Digital Twin Dev (Jr)	MEDIUM	2028–30	AI-assisted calibration, physics correlation	Model validation strategy	
	Virtual Testing Execution	HIGH	2027–29	AI auto-config of virtual test environments	Novel failure modes, safety-critical	
	FMEA & Risk Analysis (Routine)	HIGH	2026–28	AI FMEA generation from design & failure libraries	Novel failure mode discovery	
	Homologation & Type Approval	HIGH	2027–28	AI doc assembly, regulatory cross-reference	Authority relationship, novel approvals	
	Safety Case & FMEDA Docs	MEDIUM	2027–29	AI-assisted safety argumentation, FMEDA tables	Safety integrity judgment, cert sign-off	
R&D	Literature Review & SotA Survey	CRITICAL	2026–27	AI research agents (Semantic Scholar, Perplexity)	Research hypothesis framing	
	Tech Scouting & Competitor Analysis	HIGH	2026–27	AI patent landscape analysis, market intelligence	Strategic investment decisions	
	IP & Patent Research	CRITICAL	2026–27	AI prior art mining, FTO analysis tooling	IP strategy, inventor interviews	
	PoC Engineering (Jr)	MEDIUM	2028–30	AI rapid prototyping, generative design for PoC	Creative problem framing, technical risk	
	Research Data Analysis	HIGH	2026–28	AI statistical modelling, automated insight extraction	Hypothesis generation, domain expertise	
	Grant & Proposal Writing (Ops)	HIGH	2026–27	LLM grant proposal drafting, compliance structuring	Intellectual originality, funder relationships	
	Eng Programme Planning (Jr-Mid)	MEDIUM	2027–29	AI schedule generation, resource optimisation	Stakeholder alignment, scope negotiation	
	R&D Milestone Tracking	HIGH	2026–28	AI automated status reports, deviation alerts	Executive narrative, board-level decisions	
Tech Docs & KM	Technical Manual & SOP Writing	CRITICAL	2025–26	LLM end-to-end SOP generation	Safety-critical tone, regulatory approval	
	Parts & Service Docs (IETP)	CRITICAL	2025–27	AI IETP generation, illustrated parts automation	OEM review, liability sign-off	
	Training Material Dev	CRITICAL	2026–27	AI content generation, automated eLearning	Instructional design strategy	

ER&D	Clus-ter	Activity / Role	Risk	Horizon	AI Mechanism	Residual Value	Human
		Config & Interface Control Docs	HIGH	2026-27	AI document automation, change tracking	Contractual accuracy, customer review	
		Knowledge Base Curation	HIGH	2026-27	AI auto-curating, dedupli-cating, updating	Knowledge architec-ture strategy	
		Lessons Learned Capture	HIGH	2026-28	AI extraction from project data, auto-structuring	Cross-programme knowledge synthesis	

ORBYS Research

team@orbys360.com • www.orbys360.com

*This report is published for informational and strategic planning purposes only.
All rights reserved. Reproduction without written permission is prohibited.*